

Artificial Neural Networks



Mihai Oltean

www.cs.ubbcluj.ro/~moltean

moltean@cs.ubbcluj.ro

STRUCTURE

- Why we need ANNs?
 - Biological neurons and networks / brain
 - Perceptron / Artificial Neuron
 - Artificial Neural Networks
 - Training ANNs
 - Perceptron' algorithm
 - BackPropagation
 - Evolutionary Training
 - Examples
-

LEARNING PARADIGMS

- **Supervised learning**

- We have a set of example $(x, f(x))$.

- **Unsupervised learning**

- In unsupervised learning we are given some data x , and the cost function to be minimized can be any function of the data x .

- **Reinforcement learning**


- Data x is usually not given, but generated by an agent's interactions with the environment. At each point in time t , the agent performs an action y_t and the environment generates an observation x_t and an instantaneous cost c_t , according to some (usually unknown) dynamics. The aim is to discover a *policy* for selecting actions that minimizes some measure of a long-term cost, i.e. the expected cumulative cost. The environment's dynamics and the long-term cost for each policy are usually unknown, but can be estimated.
-

WHY ARTIFICIAL NEURAL NETWORKS?

- Some tasks can be done easily (effortlessly) by humans but are hard by conventional paradigms on Von Neumann machine with algorithmic approach
 - Pattern recognition (old friends, hand-written characters, voice)
 - Content addressable recall
 - Approximate, common sense reasoning (driving, playing piano, baseball player)
 - These tasks are often ill-defined, experience based, hard to apply logic
-

Human Brain

- 10.000.000.000 neurons
- 5.000 connections / neuron (in average)
- New connections between neurons can be developed during lifetime
- Small animals have fixed brain.



COMPARISON


Modern Computers

- One or a few high speed (ns) processors with considerable computing power
- One or a few shared high speed buses for communication
- Sequential memory access by address
- Problem-solving knowledge is separated from the computing component
- Hard to be adaptive

Human Brain


- Large # (10^{11}) of low speed processors (ms) with limited computing power
 - Large # (10^{15}) of low speed connections
 - Content addressable recall (CAM)
 - Problem-solving knowledge resides in the connectivity of neurons
 - Adaptation by changing the connectivity
-

• Biological neural activity



- Each neuron has a *body*, an *axon*, and many *dendrites*
- Can be in one of the two states: *firing* and *rest*.
- Neuron fires if the total incoming stimulus exceeds the threshold
- *Synapse*: thin gap between axon of one neuron and dendrite of another.
- Signal exchange
- Synaptic strength/efficiency

AN EXAMPLE OF ANN



*Schematic of a single hidden layer,
feed-forward neural network.*

WHAT IS AN (ARTIFICIAL) NEURAL NETWORK?

- A set of **nodes** (units, neurons, processing elements)
 - Each node has input and output
 - Each node performs a simple computation by its **node function**
 - **Weighted connections** between nodes
 - Connectivity gives the structure/architecture of the net
 - What can be computed by a NN is primarily determined by the connections and their weights
 - *A very much simplified version of networks of neurons in animal nerve systems*
-

ANN

Bio NN

- **Nodes**

- input
- output
- node function

- **Connections**

- connection strength

- **Cell body**

- signal from other neurons
- firing frequency
- firing mechanism

- **Synapses**

- synaptic strength
-

BENEFITS OF USING ANNs

- Learning: ANN have the ability to learn based on the so called learning stage.
 - Auto organization: a ANN creates its own representation of the data given in the learning process.
 - Tolerance to faults: because ANN store redundant information, partial destruction of the neural network do not damage completely the network response.
 - Flexibility: ANN can handle input data without important changes like noisy signals or others changes in the given input data (e.g. if the input data is an object, this can be a little different without problems to the ANN response).
 - Real Time: ANN are parallel structures; if they are implemented in this way using computers or special hardware real time can be achieved.
 - Scalability: An ANN can be easily ported to fit any problem from a particular problem area.
 - Highly parallel, simple local computation (at neuron level) achieves global results as emerging property of the interaction (at network level)
-

HISTORY OF NN

○ **Pitts & McCulloch (1943)**

- First mathematical model of biological neurons
- All Boolean operations can be implemented by these neuron-like nodes (with different threshold and excitatory/inhibitory connections).
- Competitor to Von Neumann model for general purpose computing device
- Origin of automata theory.


○ **Hebb (1949)**

- Hebbian rule of learning: increase the connection strength between neurons i and j whenever both i and j are activated.
 - Or increase the connection strength between nodes i and j whenever both nodes are simultaneously ON or OFF.
-


EARLY BOOMING (50's – EARLY 60's)

- Rosenblatt (1958)
 - Perceptron: network of threshold nodes for pattern classification
 - Perceptron learning rule
 - Perceptron convergence theorem: everything that can be represented by a perceptron can be learned
 - Widrow and Hoff (1960, 1962)
 - Learning rule based on gradient descent (with differentiable unit)
 - Minsky's attempt to build a general purpose machine with Pitts/McCulloch units
-

PERCEPTRON WITH STEP FUNCTION




Marvin Minsky (MIT AI Lab)



THE SETBACK (MID 60's – LATE 70's)

- Serious problems with perceptron model (Minsky's book 1969)
 - Single layer perceptrons cannot represent (learn) simple functions such as XOR
 - Multi-layer of non-linear units may have greater power but there is no learning rule for such nets
 - Scaling problem: connection weights may grow infinitely
 - The first two problems overcame by latter effort in 80's, but the scaling problem persists
 - Death of Rosenblatt (1964)
 - Striving of Von Neumann machine and AI
-

mlp for xor



RENEWED ENTHUSIASM AND FLOURISH (80'S – PRESENT)

- New techniques

 - Backpropagation learning for multi-layer feed forward nets (with non-linear, differentiable node functions)
 - Thermodynamic models (Hopfield net, Boltzmann machine, etc.)
 - Unsupervised learning
 - Impressive application (character recognition, speech recognition, text-to-speech transformation, process control, associative memory, etc.)
 - Traditional approaches face difficult challenges
 - Caution:
 - Don't underestimate difficulties and limitations
 - Poses more problems than solutions
-

ANN REQUIREMENTS


Structure

Activation function


Learning algorithm

Artificial Neuron Models

- Each node has one or more inputs from other nodes, and one output to other nodes
- Input/output values can be
 - Binary $\{0, 1\}$
 - Bipolar $\{-1, 1\}$
 - Continuous
- All inputs to one node come in at the same time and remain activated until the output is produced
- Weights associated with links
-



General neuron model



Weighted input summation

Node Function


○ Step (threshold) function

$$f(\text{net}) = \begin{cases} a & \text{if } \text{net} < c \\ b & \text{if } \text{net} > c \end{cases}$$


where c is called the threshold

○ Ramp function

$$f(\text{net}) = \begin{cases} a & \text{if } \text{net} \leq c \\ b & \text{if } \text{net} \geq d \\ a + \frac{(\text{net}-c)(b-a)}{(d-c)} & \text{otherwise} \end{cases}$$



Step function



Ramp function

Node Function (Sigmoid function)


- S-shaped
- Continuous and everywhere differentiable
- Rotationally symmetric about some point ($net = c$)
- Asymptotically approach saturation points

$$\lim_{net \rightarrow -\infty} f(net) = a \quad \lim_{net \rightarrow \infty} f(net) = b$$

- Examples:

$$f(net) = z + \frac{1}{1 + \exp(-x \cdot net + y)}$$

$$f(net) = \tanh(x \cdot net - y) + z,$$



Sigmoid function

When $y = 0$ and $z = 0$:

$$a = 0, b = 1, c = 0.$$

When $y = 0$ and $z = -0.5$

$$a = -0.5, b = 0.5, c = 0.$$


Larger x gives steeper curve

Node Function (Gaussian)

Bell-shaped (radial basis)

- Continuous
- $f(\text{net})$ asymptotically approaches 0 (or some constant) when $|\text{net}|$ is large
- Single maximum
- Example:


$$f(\text{net}) = \frac{1}{\sqrt{2\pi}\sigma} \exp \left[-\frac{1}{2} \left(\frac{\text{net} - \mu}{\sigma} \right)^2 \right]$$



Gaussian function

Network Architecture (Asymmetric) Fully Connected Networks

- Every node is connected to every other node
- Connection may be excitatory (positive), inhibitory (negative), or irrelevant (0).
- Most general
- Symmetric fully connected nets: weights are symmetric ($w_{ij} = w_{ji}$)



Input nodes:
receive input from the environment


Output nodes: send signals to the environment

Hidden nodes: no direct interaction to the environment

Network Architecture

Layered Networks

- Nodes are partitioned into subsets, called layers.
- No connections that lead from nodes in layer j to those in layer k if $j > k$.




- Inputs from the environment are applied to nodes in layer 0 (**input layer**).
- Nodes in input layer are placeholders with no computation occurring (i.e., their node functions are identity function)

Network Architecture

Feedforward Networks

- A connection is allowed from a node in layer i only to nodes in layer $i + 1$.
- Most widely used architecture.



Conceptually, nodes at higher levels successively abstract features from preceding layers

What kind of problems can ANNs solve?

- **Regression (function approximation)**
 - **Classification**
-

TRAINING AN ANN

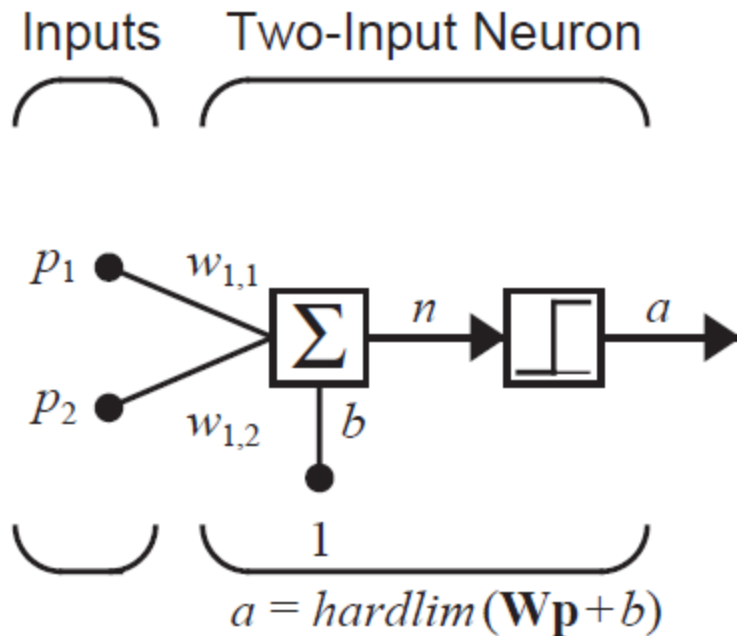
- Perceptron' algorithm
 - Back-propagation
 - Evolutionary
 -
-

PERCEPTRON' ALGORITHM


- A linear classifier for classifying data X
 - Parameters are adapted with an ad-hoc rule similar to stochastic steepest gradient descent.
 - Can only perfectly classify a set of data for which different classes are linearly separable in the input space,
 - It often fails completely for non-separable data.
-

EXAMPLE

TRAINING A SIMPLE PERCEPTRON FOR A CLASSIFICATION PROBLEM



NODE FUNCTION




Hard-Limit Transfer Function

The perceptron produces:

=1 if the net input into the transfer function is equal to or greater than 0;

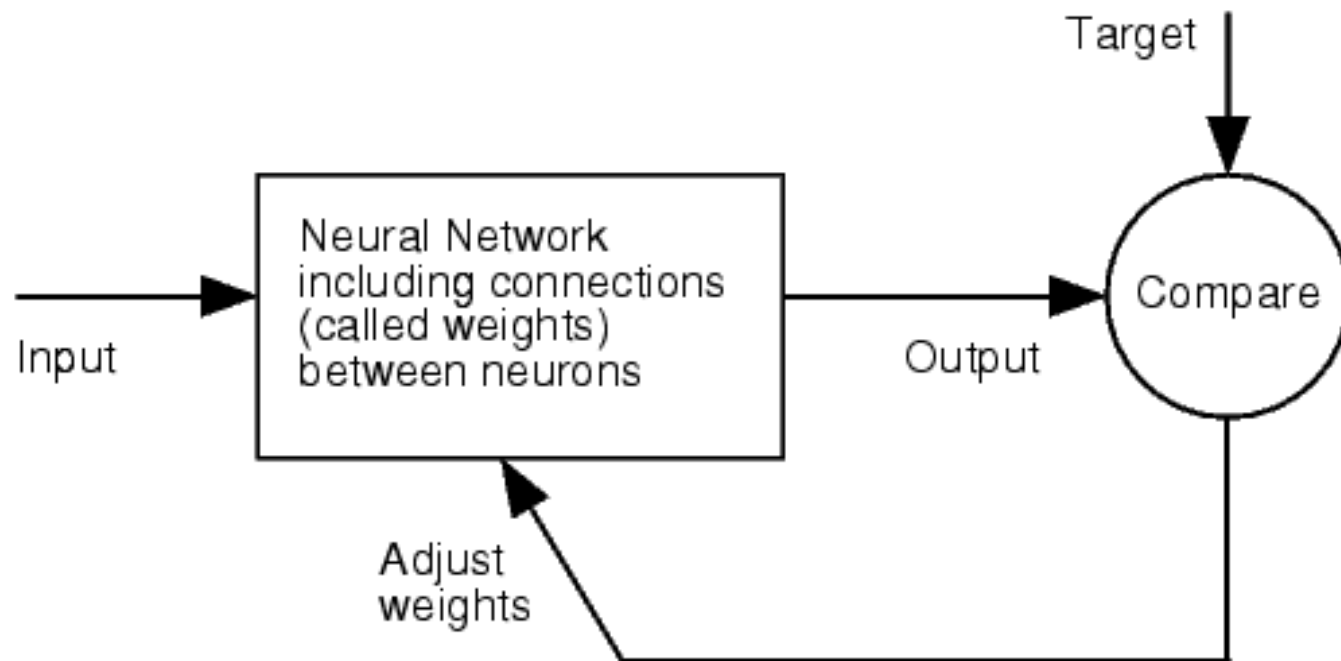
=otherwise it produces a 0.

DECISION BOUNDARY



$$\begin{aligned}w_1 &= 1 \\w_2 &= 1 \\b &= -1\end{aligned}$$

TRAINING ALGORITHM



RULES FOR TRAINING (2)





- If an input vector is presented and the output of the neuron is correct ($\mathbf{a} = \mathbf{t}$, and $\mathbf{e} = \mathbf{t} - \mathbf{a} = 0$), then the weight vector \mathbf{w} is not altered.
- If the neuron output is 0 and should have been 1 ($\mathbf{a} = 0$ and $\mathbf{t} = 1$, and $\mathbf{e} = \mathbf{t} - \mathbf{a} = 1$), the input vector \mathbf{p} is added to the weight vector \mathbf{w} .

This makes the weight vector point closer to the input vector, increasing the chance that the input vector will be classified as a 1 in the future.

- If the neuron output is 1 and should have been 0 ($\mathbf{a} = 1$ and $\mathbf{t} = 0$, and $\mathbf{e} = \mathbf{t} - \mathbf{a} = -1$), the input vector \mathbf{p} is subtracted from the weight vector \mathbf{w} .

This makes the weight vector point farther away from the input vector, increasing the chance that the input vector is classified as a 0 in the future.

RULES FOR TRAINING



TRAINING DATA

$$\left\{ \mathbf{p}_1 = \begin{bmatrix} 2 \\ 2 \end{bmatrix}, t_1 = 0 \right\} \quad \left\{ \mathbf{p}_2 = \begin{bmatrix} 1 \\ -2 \end{bmatrix}, t_2 = 1 \right\} \quad \left\{ \mathbf{p}_3 = \begin{bmatrix} -2 \\ 2 \end{bmatrix}, t_3 = 0 \right\} \quad \left\{ \mathbf{p}_4 = \begin{bmatrix} -1 \\ 1 \end{bmatrix}, t_4 = 1 \right\}$$

First step

$$\mathbf{W}(0) = \begin{bmatrix} 0 & 0 \end{bmatrix} \quad b(0) = 0$$

$$a = \text{hardlim}(\mathbf{W}(0)\mathbf{p}_1 + b(0))$$

$$= \text{hardlim}\left(\begin{bmatrix} 0 & 0 \end{bmatrix} \begin{bmatrix} 2 \\ 2 \end{bmatrix} + 0\right) = \text{hardlim}(0) = 1$$

$$e = t_1 - a = 0 - 1 = -1$$

$$\Delta \mathbf{W} = e \mathbf{p}_1^T = (-1) \begin{bmatrix} 2 & 2 \end{bmatrix} = \begin{bmatrix} -2 & -2 \end{bmatrix}$$

$$\Delta b = e = (-1) = -1$$

$$\mathbf{W}^{new} = \mathbf{W}^{old} + e \mathbf{p}^T = \begin{bmatrix} 0 & 0 \end{bmatrix} + \begin{bmatrix} -2 & -2 \end{bmatrix} = \begin{bmatrix} -2 & -2 \end{bmatrix} = \mathbf{W}(1)$$

$$b^{new} = b^{old} + e = 0 + (-1) = -1 = b(1)$$

SECOND STEP

$$\begin{aligned} a &= \text{hardlim}(\mathbf{W}(1)\mathbf{p}_2 + b(1)) \\ &= \text{hardlim}\left(\begin{bmatrix} -2 & -2 \end{bmatrix} \begin{bmatrix} 1 \\ -2 \end{bmatrix} - 1\right) = \text{hardlim}(1) = 1 \end{aligned}$$

Ok, no modification to the weights

THE SOLUTION

$$\mathbf{w}(6) = \begin{bmatrix} -2 & -3 \end{bmatrix}$$

$$b(6) = 1$$

BACK PROPAGATION - WHY

- Networks without hidden units are very limited in the input-output mappings they can model.
 - More layers of linear units do not help. Its still linear.
 - Fixed output non-linearities are not enough
 - We need multiple layers of adaptive non-linear hidden units. This gives us a universal approximator. But how can we train such nets?
 - We need an efficient way of adapting **all** the weights, not just the last layer. This is hard. Learning the weights going into hidden units is equivalent to learning features.
 - Nobody is telling us directly what hidden units should do.
-

BACK PROPAGATION ALGORITHM

Initialize all weights to small random numbers.

Repeat

For each example X do



- Propagate example X forward through the network

- Propagate errors backward through the network

Until error is small

THE IDEA BEHIND BACKPROPAGATION

- We don't know what the hidden units ought to do, but we can compute how fast the error changes as we change a hidden activity.
 - Instead of using desired activities to train the hidden units, use **error derivatives w.r.t. hidden activities**.
 - Each hidden activity can affect many output units and can therefore have many separate effects on the error. These effects must be combined.
 - We can compute error derivatives for **all** the hidden units efficiently.
 - Once we have the error derivatives for the hidden activities, it's easy to get the error derivatives for the weights going into a hidden unit.
-



ADVANTAGES AND DISADVANTAGES OF BP

- Reduced running time – smaller than evolutionary training (Strength)
 - Cannot train activation functions that have no derivative – (Weak)
 - Can easily fall into local optima – (Weak)
 - Cannot discover new connections between neurons – (Weak)
-

EVOLUTIONARY TRAINING

- A solution is a set of values for the weight of each connection between two neurons.
 - Start with a set of random weights.
 - Compute the quality of that NN.
 - Repeat
 - Mutation / Crossover / Fitness
 - Until some good values have been obtained
-

ADVANTAGES AND DISADVANTAGES

- Big running time – bigger than back-propagation (Weak)
 - Can train any kind of activation function (including those that have no derivate) – (Strength)
 - Can escape of local optima better than backpropagation – (Strength)
 - Can discover new connections between neurons – (Strength)
-

DIFFERENCES BETWEEN ANNs AND GP

- GP discover the entire structure

ANNs discover only the weight.

- Connection between nodes have weight 1

Connections between nodes have a weight

- Within a node we can have any function

Within a node we have a predefined node function

- A node always fires a value

A neuron does not always fires values.

SOLVING PROBLEMS WITH ANNs

- Regression
 - Classification
 - (training similar with GP)
 - Some fitness cases are needed
-

ANNs FOR DIGITS RECOGNITION

- Each digit is represented over a matrix of pixels (no more than 8x8)


- ANN with 64 inputs, 1 output and multiple hidden layers.

- Output between

- $[a_1, a_2) \rightarrow 0$
- $[a_2, a_3) \rightarrow 1$
- ...
- $[a_{10}, a_{11}) \rightarrow 9$

- 10 training data

- $(00010101, 0)$
- $(10101010, 1)$
- ...
- $(10101011, 9)$



REFERENCES

- Martin T. Hagan, Howard B. Demuth, Mark H. Beale, Neural Network Design, PWS Publishing, 1996
-