

Computer Networks

The Network Layer TCP/IP

Adrian Sergiu DARABANT

Lecture
8

IP Datagram

IP protocol version
number

header length
(bytes)

"type" of data

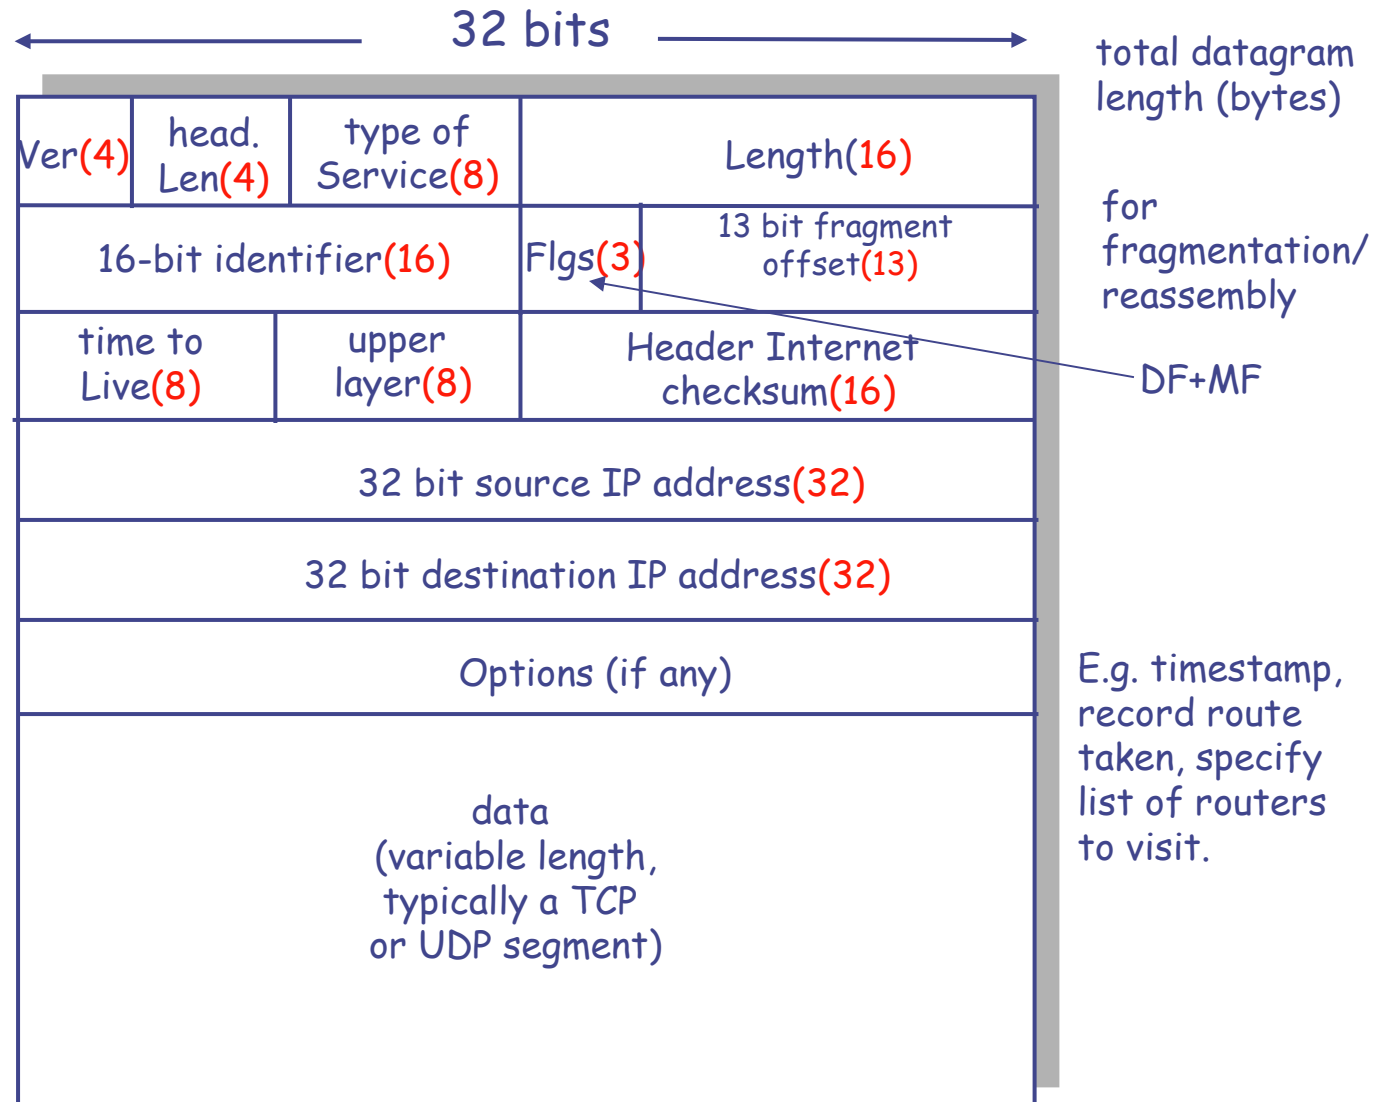
max number
remaining hops
(decremented at
each router)

upper layer protocol
to deliver payload to

how much

overhead with
IP?

- ◆ 20 bytes of IP
- ◆ + transp layer overhead



Datagram: from source to destination



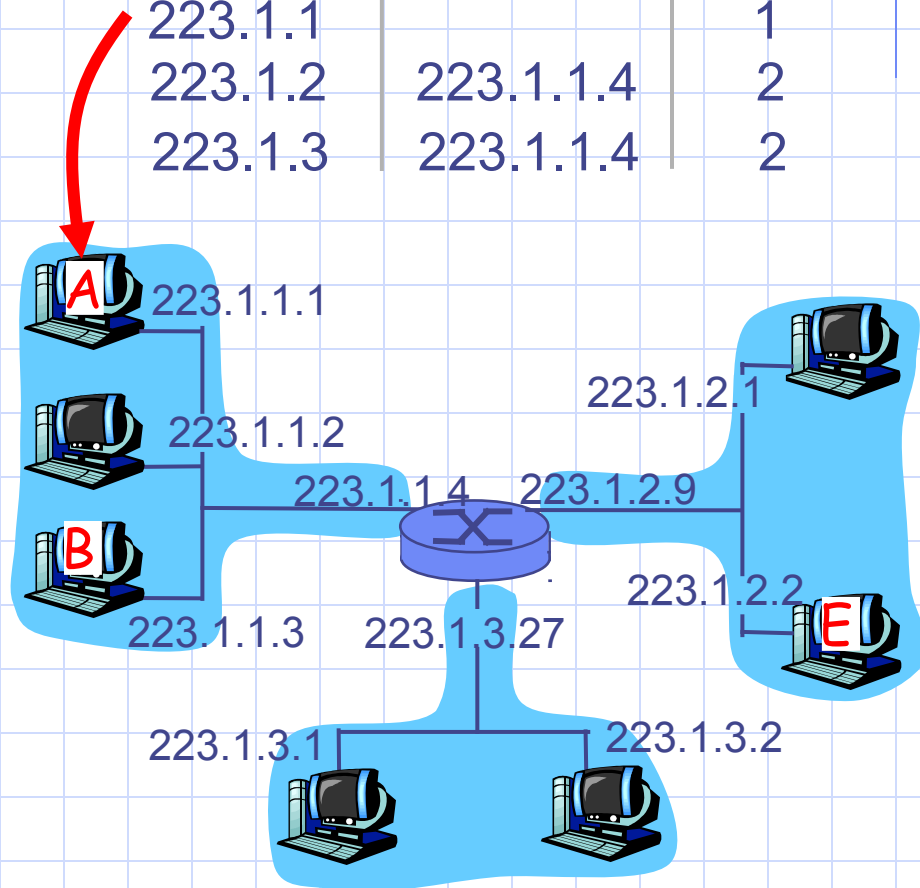
IP datagram:

misc fields	source IP addr	dest IP addr	data
----------------	-------------------	-----------------	------

- ◆ datagram remains **unchanged**, as it travels source to destination
- ◆ addr fields of interest here

forwarding table in

Dest. Net.	next router	Nhops
223.1.1		1
223.1.2	223.1.1.4	2
223.1.3	223.1.1.4	2



Datagram: from source to destination

misc fields	223.1.1.1	223.1.1.3	data
-------------	-----------	-----------	------

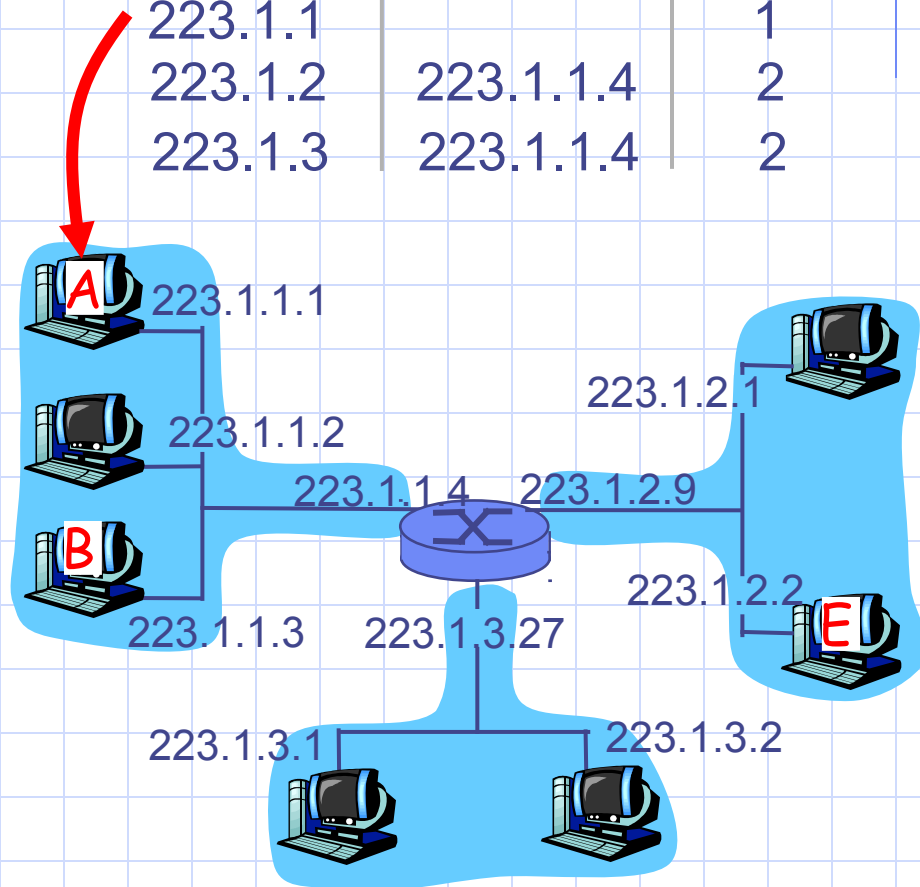
Starting at A, send IP datagram addressed to B:

- ◆ look up net. address of B in forwarding table
- ◆ find B is on same net. as A
- ◆ link layer will send datagram directly to B inside link-layer frame
 - B and A are directly connected

forwarding table in

Dest. Net.	next router	Nhops
------------	-------------	-------

223.1.1		1
223.1.2	223.1.1.4	2
223.1.3	223.1.1.4	2



Datagram: from source to destination

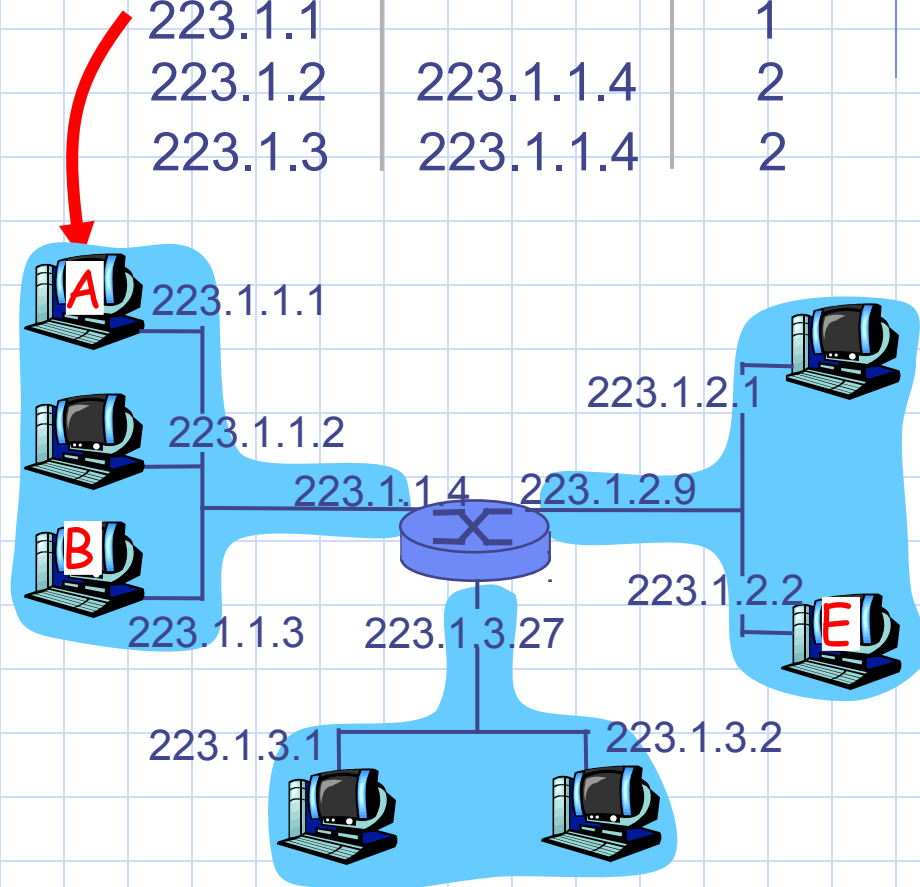
misc fields	223.1.1.1	223.1.2.3	data
-------------	-----------	-----------	------

Starting at A, dest. E:

- ◆ look up network address of E in forwarding table
- ◆ E on *different* network
 - A, E not directly attached
- ◆ routing table: next hop router to E is 223.1.1.4
- ◆ link layer sends datagram to router 223.1.1.4 inside link-layer frame
- ◆ datagram arrives at 223.1.1.4

forwarding table in

Dest. Net.	next router	Nhops
223.1.1		1
223.1.2	223.1.1.4	2
223.1.3	223.1.1.4	2



Datagram: from source to destination


misc fields	223.1.1.1	223.1.2.3	data
-------------	-----------	-----------	------

Arriving at 223.1.4,
destined for 223.1.2.2

- ◆ look up network address of E in router's forwarding table
- ◆ E on *same* network as router's interface 223.1.2.9
 - router, E directly attached
- ◆ link layer sends datagram to 223.1.2.2 inside link-layer frame via interface 223.1.2.9


forwarding table in

Dest. Network	router	Nhops	interface
223.1.1	-	1	223.1.1.4
223.1.2	-	1	223.1.2.9
223.1.3	-	1	223.1.3.27



Fragmentation/Reassembly

- ◆ network links have MTU (max.transfer size) - largest possible link-level frame.
 - different link types, different MTUs
- ◆ large IP datagram divided (“fragmented”) within net
 - one datagram becomes several datagrams
 - “reassembled” only at final destination
 - IP header bits used to identify, order related



Fragmentation/Reassembly

y

Example

- ◆ 4000 byte datagram
- ◆ MTU = 1500 bytes

length	ID	fragflag	offset
=4000	=x	=0	=0

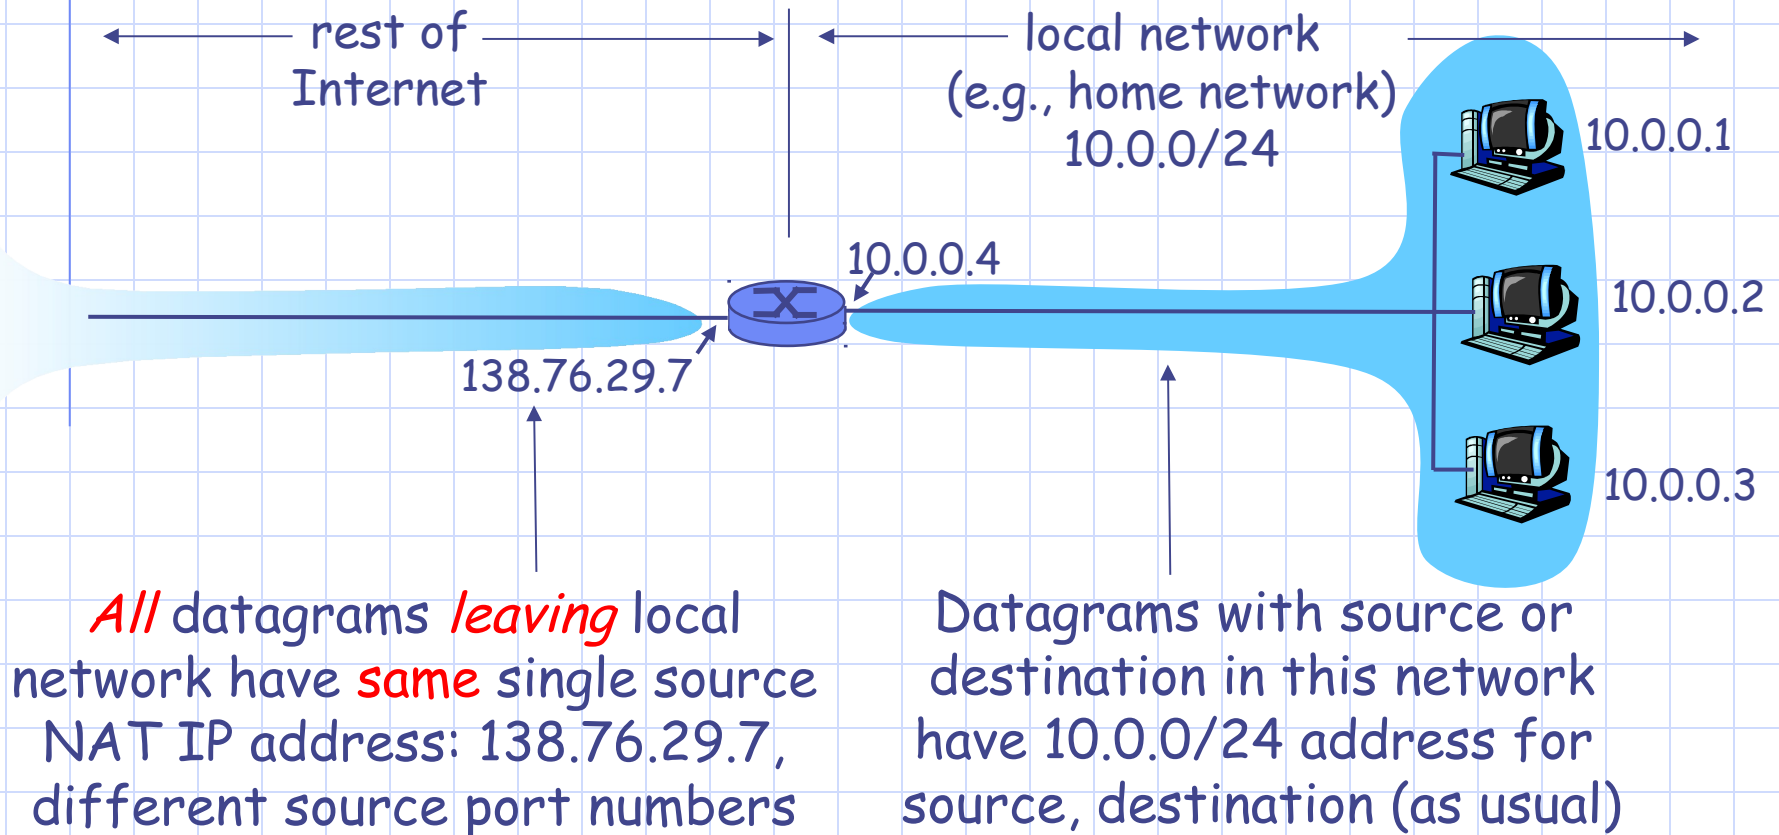
One large datagram becomes several smaller datagrams

length	ID	fragflag	offset
=1500	=x	=1	=0

length	ID	fragflag	offset
=1500	=x	=1	=1480

length	ID	fragflag	offset
=1040	=x	=0	=2960


NAT – Network Address Translation



NAT – Network Address Translation

- ◆ **Motivation:** local network uses just one IP address as far as outside world is concerned:
 - no need to be allocated range of addresses from ISP: - just one IP address is used for all devices
 - can change addresses of devices in local network without notifying outside world
 - can change ISP without changing addresses of devices in local network
 - devices inside local net not explicitly addressable, visible by outside world (a security plus)

NAT – Network Address Translation



NAT – Network Address Translation

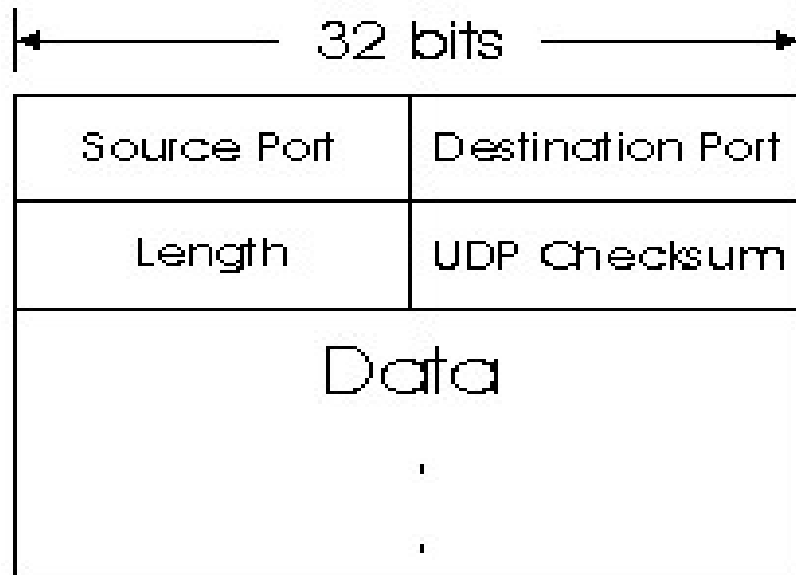
◆ 16-bit port-number field:

- 60,000 simultaneous connections with a single LAN-side address!

◆ NAT is controversial:

- routers should only process up to layer 3
- violates end-to-end argument
 - ◆ NAT possibility must be taken into account by app designers, e.g., P2P applications
- address shortage should instead be

UDP



how much overhead with UDP?

- ◆ 20 bytes of IP
- ◆ 8 bytes of UDP
- ◆ = 28 bytes + app layer overhead

Checksum – for the entire datagram (header + data)

Length ≥ 8 – entire datagram

ICMP

- ◆ Used by hosts, routers, gateways to communication network-level information
 - error reporting: unreachable host, network, port, protocol
 - echo request/reply (used by ping)
- ◆ Network-layer “above” IP:
 - ICMP msgs carried in IP datagrams
- ◆ **ICMP message:** type, code plus first 8 bytes of IP datagram causing error

UDP Rules

Unreliable – When a message is sent, it cannot be known if it will reach its destination; it could get lost along the way. There is no concept of acknowledgment, retransmission, or timeout.

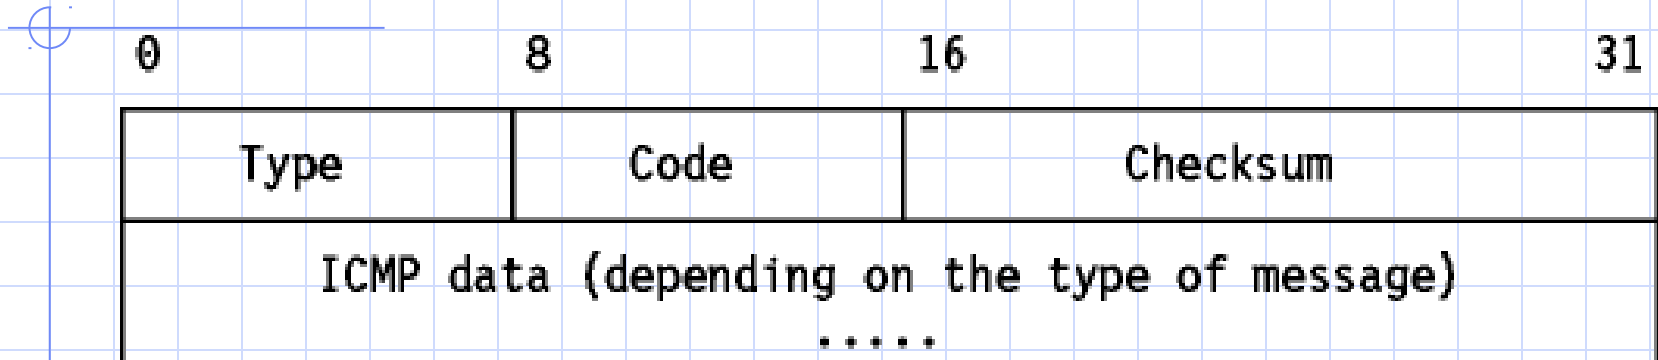
Not ordered – If two messages are sent to the same recipient, the order in which they arrive cannot be predicted.

Lightweight – There is no ordering of messages, no tracking connections, etc. It is a small transport layer designed on top of IP.

Datagrams – Packets are sent individually and are checked for integrity only if they arrive. Packets have definite boundaries which are honored upon receipt, meaning a read operation at the receiver socket will yield an entire message as it was originally sent.

No congestion control – UDP itself does not avoid congestion, and it's possible for high bandwidth applications to trigger congestion collapse, unless they implement

ICMP




<u>Type</u>	<u>Code</u>	<u>description</u>	<u>Type</u>	<u>Code</u>	<u>description</u>
0	0	echo reply (ping)	4	0	source quench (congestion control - not used)
3	0	dest. network unreachable	8	0	echo request (ping)
3	1	dest host unreachable	9	0	route advertisement
3	2	dest protocol unreachable	10	0	router discovery
3	3	dest port unreachable	11	0	TTL expired
3	6	dest network unknown	12	0	bad IP header
3	7	dest host unknown			

Network diagnostic

- ◆ **Ping** – uses ICMP **Echo** and **Reply** to determine if a host is “up”
- ◆ **Traceroute** – determine the path (as routers) from a source host to a destination host using UDP(usually).

TCP Datagrams



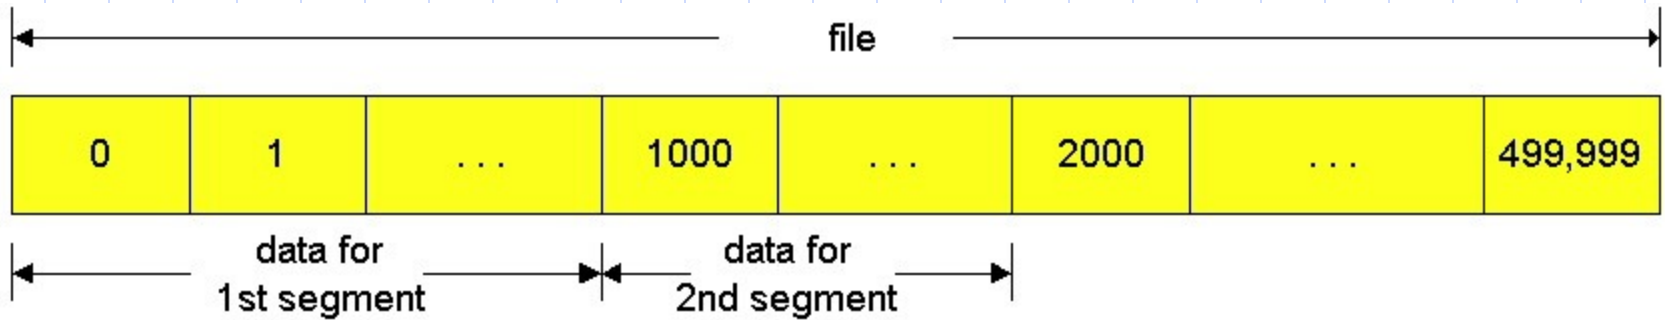
how much
overhead with
TCP?

- ◆ 20 bytes of TCP
- ◆ 20 bytes of IP
- ◆ = 40 bytes + app layer overhead

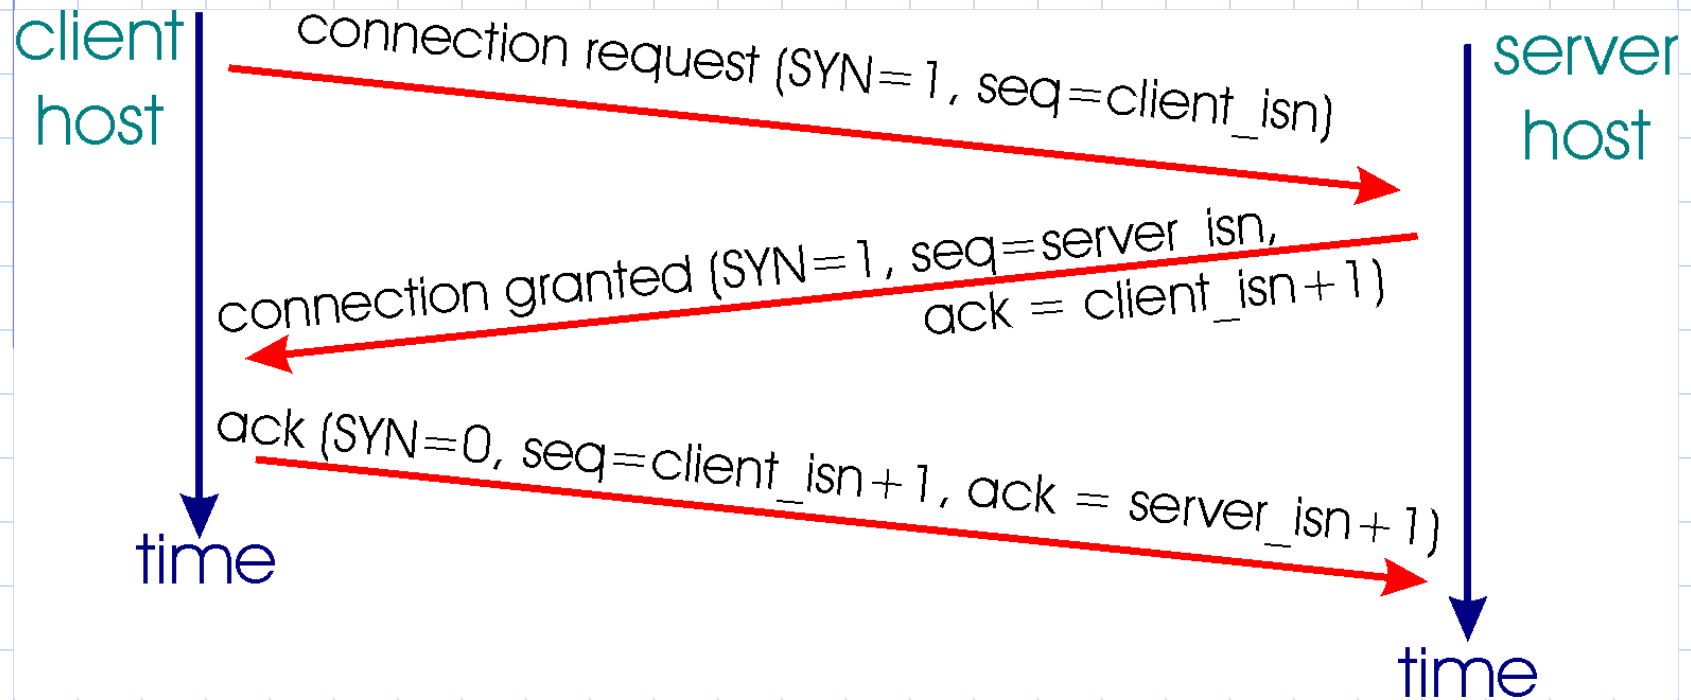
TCP - Data Transfer

- **Ordered data transfer** — the destination host rearranges according to sequence number
- **Retransmission of lost packets** — any cumulative stream not acknowledged is retransmitted
- **Error-free data transfer**
- **Flow control** — limits the rate a sender transfers data to guarantee reliable delivery. The receiver continually hints the sender on how much data can be received (controlled by the sliding window). When the receiving host's buffer fills, the next acknowledgment contains a 0 in the window size, to stop transfer and allow the data in the buffer to be processed.
- **Congestion control**

TCP Segments



TCP Open - 3-way handshake



TCP Connection Teardown


Closing a connection:

client closes socket:

```
clientSocket.close();
```

Step 1: client end system
sends TCP FIN control
segment to server_

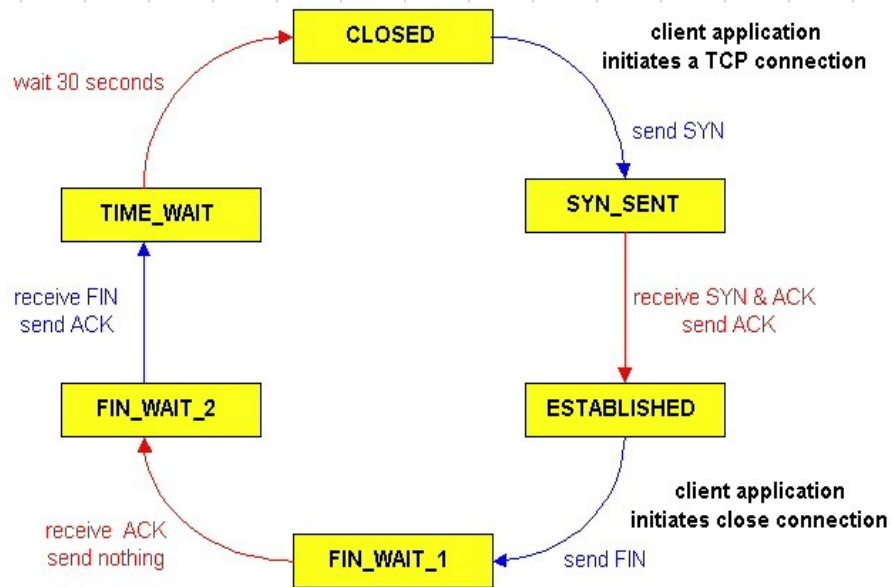
Step 2: server receives
FIN, replies with ACK.
Closes connection, sends
FIN.



Seq numbers and Acks

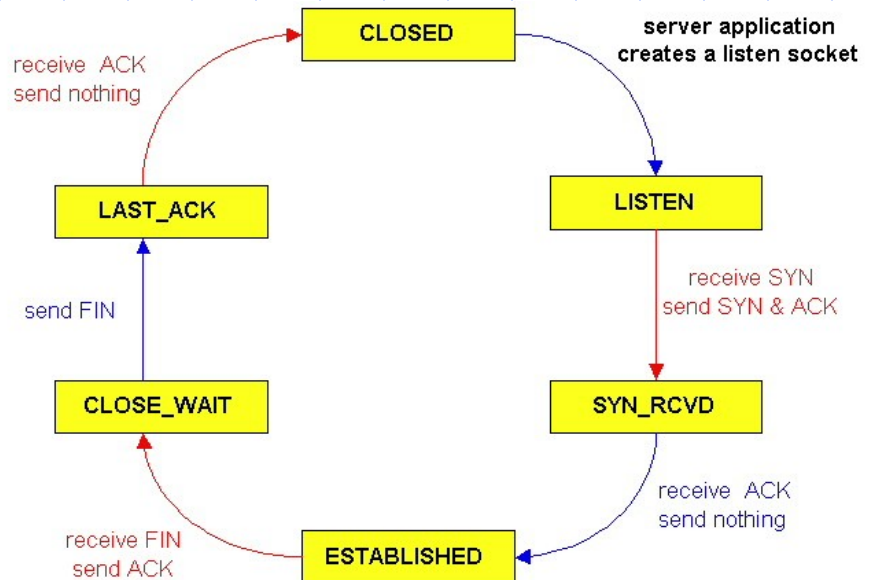
- ◆ Sequence numbers are used to reassemble data in the order in which it was sent.
- ◆ Sequence numbers increment based on the number of bytes in the TCP data field.
 - -Known as a Byte Sequencing Protocol
- ◆ Each segment transmitted must be acknowledged.
 - -Multiple segments can be acknowledged
- ◆ The ACK (Acknowledgement) field indicates the next byte (sequence) number the receiver expects to receive.
- ◆ The sender, no matter how many transmitted segments, expects to receive an ACK that is one more than the number of the last transmitted byte.

TCP States



TCP
client

TCP Server




TCP Flow & Window Control

- ◆ **Sliding Window mechanism** -> the number of allowed unacknowledged bytes
 - Stop and Wait
 - Go-back N (TCP)
 - Selective Repeat
- ◆ Receiver Window
- ◆ Sender Window


Go Back N

If seg k not received \Rightarrow discard $k+1, k+2, \text{ etc}$



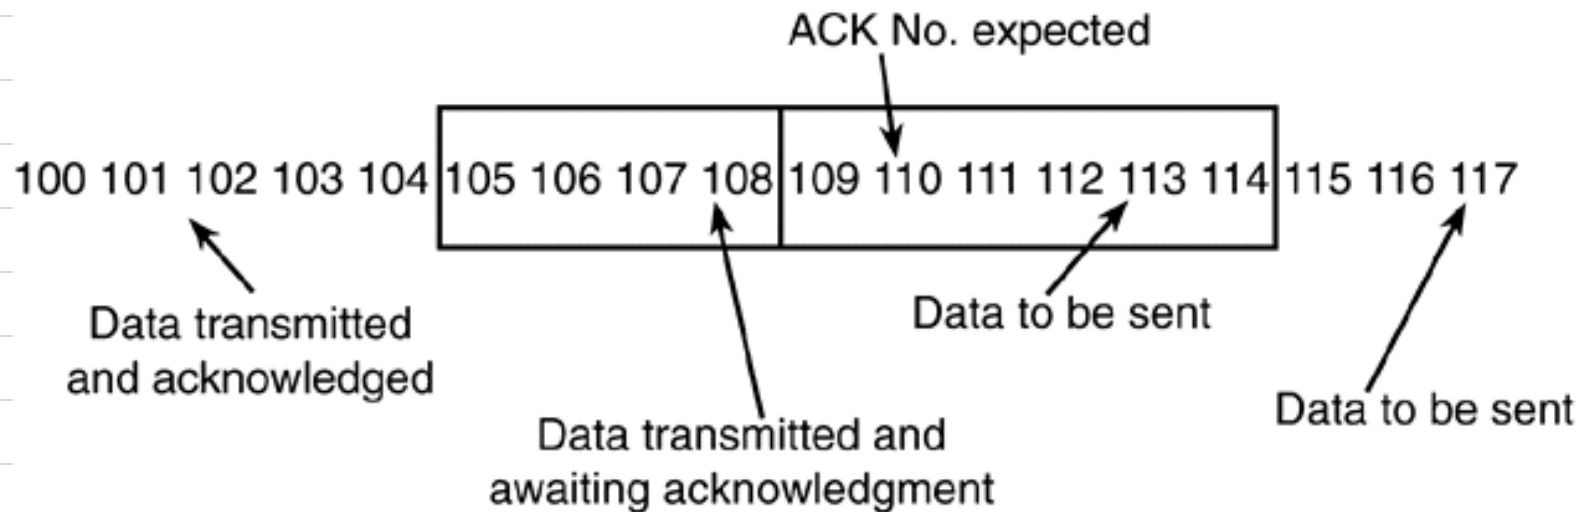
This implicitly sets the Window Size = 1

Selective Repeat



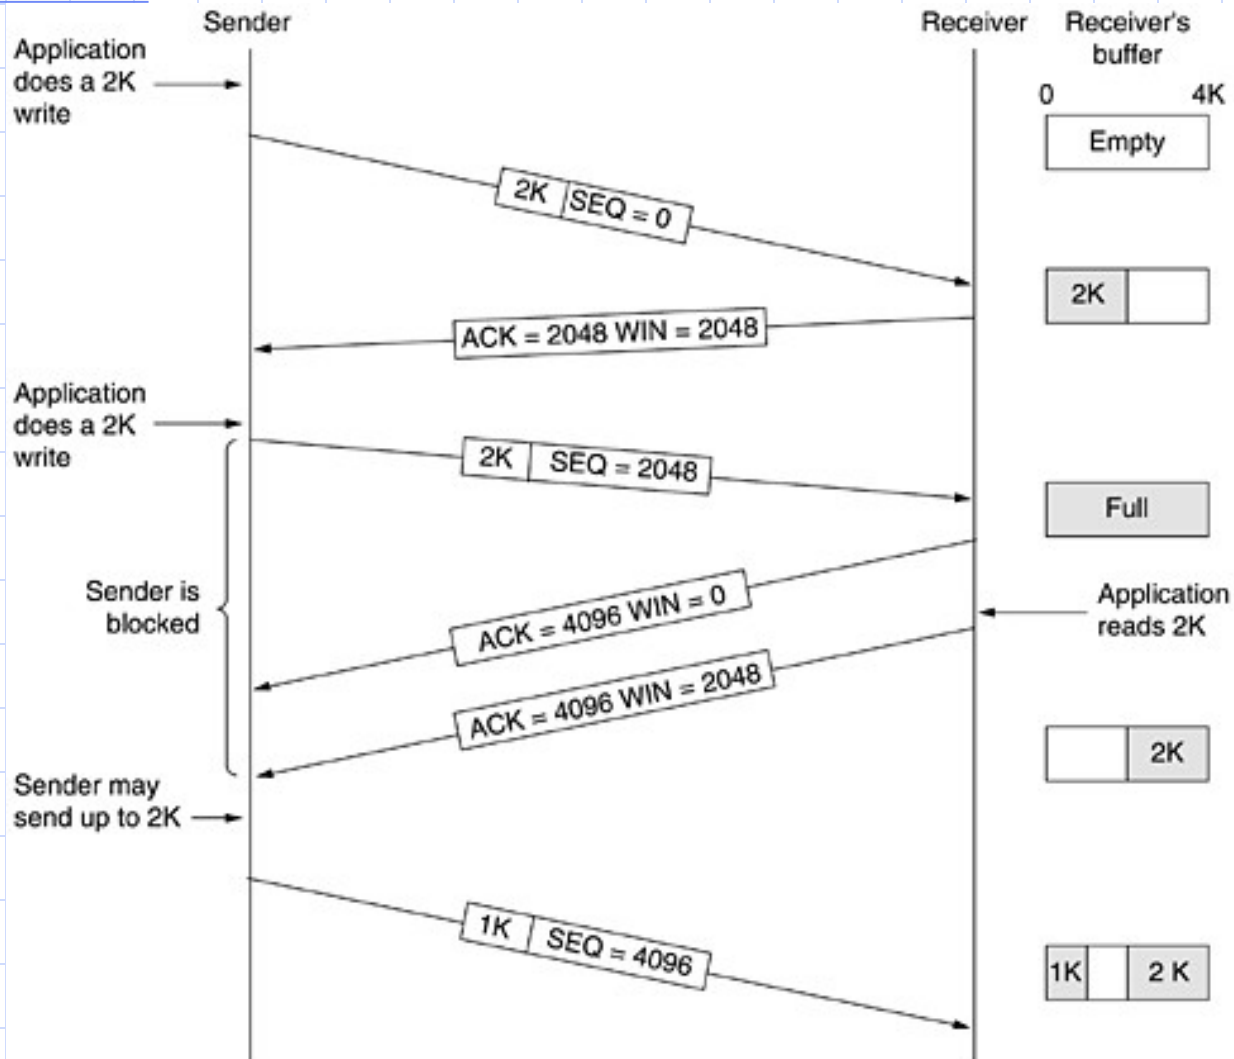
TCP Send Window

Windows based on advertised window in the received packet from the partner



Note: The actual segment size is usually 512 or 536 bytes each, but for clarity, I have shown a much smaller size.

Window Management



TCP Retransmission

- ◆ TCP will retransmit a segment upon expiration of an adaptive transmission timer.
- ◆ The timer is variable.
- ◆ When TCP transmits a segment, it records the time of transmission and the sequence number of the segment.
- ◆ When TCP receives an acknowledgment, it records the time.
- ◆ This allows TCP to build a sample round-trip delay time. **(RTT)**
- ◆ TCP will build an average delay time for a packet to be sent and received.
- ◆ The timer is slowly changed to allow for the varying differences in the Internet.


Timeout value ?

$$\text{EstimatedRTT} = (1 - \alpha) \text{EstimatedRTT} + \alpha \text{ SampleRTT}$$

$$\text{DevRTT} = (1 - \beta) \text{DevRTT} + \beta | \text{SampleRTT} - \text{EstimatedRTT} |$$


$$\text{TimeoutInterval} = \text{EstimatedRTT} + 4 \text{ DevRTT}$$

Retransmission-1




Retransmission-2

Host
A



Retransmission-3



Principles of Congestion Control

Congestion:

- ◆ informally: “too many sources sending too much data too fast for *network* to handle”
- ◆ different from flow control!
- ◆ manifestations:
 - lost packets (buffer overflow at routers)
 - long delays (queueing in router buffers)
- ◆ a top-10 problem!

Approaches towards congestion control

End-end congestion control:

- ✂ no explicit feedback from network
- ✂ congestion inferred from end-system observed loss, delay
- ✂ approach taken by TCP

Network-assisted congestion control:

- ◆ routers provide feedback to end systems
 - single bit indicating congestion (SNA, DECbit, TCP/IP ECN, ATM)
 - explicit rate sender should send at

Congestion Control

- ◆ Previously, TCP would start to transmit as much data as was allowed in the advertised window.
- ◆ What about congestion ? What is it ?
- ◆ A new window was added called the **congestion window**. –It is not negotiated, it is assumed. It starts out with *one segment* !

TCP Congestion Control

◆ end-end control (no network assistance)

◆ sender limits transmission:

$$\text{LastByteSent} - \text{LastByteAcked} \leq \text{CongWin}$$

◆ Roughly,

$$\text{rate} = \frac{\text{CongWin}}{\text{RTT}} \text{ Bytes/sec}$$

◆ CongWin is **dynamic**,
function of perceived network congestion

How does sender perceive congestion?

◆ loss event = timeout or 3 duplicate acks

◆ TCP sender reduces rate (CongWin) after loss event
three mechanisms:

- AIMD (*additive increase, multiplicative decrease*)
- slow start
- conservative after timeout events


TCP Slow Start

- ◆ **When connection begins, CongWin = 1 MSS**
 - Example: MSS = 500 bytes & RTT = 200 msec
 - initial rate = 20 kbps
- ◆ **available bandwidth may be \gg MSS/RTT**
 - desirable to quickly ramp up to respectable rate

When connection begins, increase rate exponentially fast until first loss event

TCP Slow Start -2

- ◆ When connection begins, increase rate exponentially until first loss event:
 - double **CongWin** every RTT
 - done by incrementing **CongWin** for every ACK received
- ◆ Summary: initial rate is slow but ramps up exponentially fast



Refinement

- ◆ After 3 dup ACKs:
 - CongWin is cut in half
 - window then grows linearly
- ◆ But after timeout event:
 - CongWin instead set to 1 MSS;
 - window then grows exponentially
 - to a threshold, then grows linearly

Philosophy:

- 3 dup ACKs indicates network capable of delivering some segments
- timeout before 3 dup ACKs is “more alarming”


Refinement -2

Q: When should the exponential increase switch to linear?

A: When CongWin gets to 1/2 of its value before timeout.

Implementation:

- ◆ Variable Threshold
- ◆ At loss event, Threshold is set to 1/2 of CongWin just before loss event




Summary: TCP Congestion Control

- ◆ When **CongWin** is below **Threshold**, sender in **slow-start** phase, window grows exponentially.
- ◆ When **CongWin** is above **Threshold**, sender is in **congestion-avoidance** phase, window grows linearly.
- ◆ When a **triple duplicate ACK** occurs,
 - $\text{Threshold} = \text{CongWin} / 2$
 - $\text{CongWin} = \text{Threshold}$.
- ◆ When **timeout** occurs,
 - $\text{Threshold} = \text{CongWin} / 2$
 - $\text{CongWin} = 1 \text{ MSS}$

TCP Fairness


Fairness goal: if K TCP sessions share same bottleneck link of bandwidth R , each should have average rate of R/K



Why is TCP fair ?

Two competing sessions:

- ◆ Additive increase gives slope of 1, as throughput increases
- ◆ multiplicative decrease decreases throughput proportionally



Fairness !!!!

Fairness and UDP

- ◆ Multimedia apps often do not use TCP
 - do not want rate throttled by congestion control
- ◆ Instead use UDP:
 - pump audio/video at constant rate, tolerate packet loss
- ◆ Research area: TCP friendly

Fairness and parallel TCP connections

- ◆ nothing prevents app from opening parallel cncctions between 2 hosts.
- ◆ Web browsers do this
- ◆ Example: link of rate R supporting 9 cncctions;
 - new app asks for 1 TCP, gets rate $R/10$
 - **new app asks for 11 TCPs, gets $R/2$!**