

Computer Networks

Routing

Adrian Sergiu DARABANT

Lecture
9


Routing

Routing protocol

Goal: determine “good” path (sequence of routers) thru network from source to dest.

Graph abstraction for routing algorithms:

- ◆ graph nodes are routers
- ◆ graph edges are physical links
 - link cost: delay, \$ cost, or congestion level



- ◆ “good” path:
 - typically means minimum cost path
 - other def’s possible

Routing Algorithm

Global or decentralized
information?

Global:

- ◆ all routers have complete topology, link cost info
- ◆ “link state” algorithms

Decentralized:

- ◆ router knows physically-connected neighbors, link costs to neighbors
- ◆ iterative process of computation, exchange of info with neighbors
- ◆ “distance vector” algorithms

Static or dynamic?

Static:

- ◆ routes change slowly over time

Dynamic:

- ◆ routes change more quickly
 - periodic update
 - in response to link cost changes

Routing tables - Campus

Destination	Gateway	Genmask	Flags	Metric	Iface
193.226.40.128	*	255.255.255.224	U	0	eth1
192.168.1.0	172.30.5.19	255.255.255.0	UG	0	eth1
192.168.0.0	172.30.1.4	255.255.255.0	UG	0	eth1
193.231.20.0	*	255.255.255.0	U	0	eth0
172.30.0.0	*	255.255.0.0	U	0	eth1
169.254.0.0	*	255.255.0.0	U	0	eth1
127.0.0.0	*	255.0.0.0	U	0	lo
default	193.231.20.9	0.0.0.0	UG	0	eth0

Routing tables (static)

Destination	Gateway	Genmask	Flags	Metric	Ref	Use	Iface
172.16.25.1	172.30.0.4	255.255.255.255	UGH	0	0	0	Eth1
193.226.40.128	0.0.0.0	255.255.255.224	U	0	0		Eth0
193.0.225.0	0.0.0.0	255.255.255.0	U	0	0		Eth0
193.231.20.0	0.0.0.0	255.255.255.0	U	0	0		Eth0
172.30.0.0	0.0.0.0	255.255.0.0	U	0	0		Eth1
169.254.0.0	0.0.0.0	255.255.0.0	U	0	0		Eth1
0.0.0.0	193.0.225.9	0.0.0.0	UG	0	0		Eth0

The **route** command – (Windows/Linux/other OS)

A Link-State Routing Algorithm

Dijkstra's algorithm

- ◆ net topology, link costs known to all nodes
 - accomplished via “link state broadcast”
 - all nodes have same info
- ◆ computes least cost paths from one node (“source”) to all other nodes
 - gives routing table for that node
- ◆ iterative: after k iterations, know least cost path to k dest.'s

Notation:

- ◆ $c(i,j)$: link cost from node i to j . cost infinite if not direct neighbors
- ◆ $D(v)$: current value of cost of path from source to dest. v
- ◆ $p(v)$: predecessor node along path from source to v , that is next v
- ◆ N : set of nodes whose least cost path definitively known

Dijkstra's Algorithm

1 **Initialization:**

2 $N = \{A\}$

3 for all nodes v

4 if v adjacent to A

5 then $D(v) = c(A, v)$

6 else $D(v) = \text{infinity}$

7

8 **Loop**

9 find w not in N such that $D(w)$ is a minimum

10 add w to N

11 update $D(v)$ for all v adjacent to w and not in N :

12 $D(v) = \min(D(v), D(w) + c(w, v))$


13 /* new cost to v is either old cost to v or known

14 shortest path cost to w plus cost from w to v */

15 **until all nodes in N**

Dijkstra's algorithm: example

Step	start N	D(B),p(B)	D(C),p(C)	D(D),p(D)	D(E),p(E)	D(F),p(F)
→ 0	A	2,A	5,A	1,A	infinity	infinity
→ 1	AD	2,A	4,D		2,D	infinity
→ 2	ADE	2,A	3,E			4,E
→ 3	ADEB		3,E			4,E
→ 4	ADEBC					4,E
5	ADEBCF					



Dijkstra's algorithm, discussion

Algorithm complexity: n nodes


Each iteration: need to check all nodes, w, not in N

$n(n+1)/2$ comparisons: $O(n^2)$


more efficient implementations possible: $O(n \log n)$

Oscillations possible:


e.g., link cost = amount of carried traffic




initially



... recompute
routing



... recompute



... recompute

Distance Vector Routing Algorithm

iterative:

- ◆ continues until no nodes exchange info.
- ◆ *self-terminating*: no “signal” to stop

asynchronous:

- ◆ nodes need *not* exchange info/iterate in lock step!

distributed:


- ◆ each node communicates *only* with directly-attached neighbors

Distance Table data structure

- ◆ each node has its own
- ◆ row for each possible destination
- ◆ column for each directly-attached neighbor to node
- ◆ example: in node X, for dest. Y via neighbor Z:

$$\begin{aligned} D^X(Y,Z) &= \text{distance from X to Y, via Z as next hop} \\ &= c(X,Z) + \min_w \{D^Z(Y,w)\} \end{aligned}$$

Distance Table: example



cost to destination via

$D^E()$	A	B	D
A	1	14	5
B	7	8	5
C	6	9	4
D	4	11	2

destination

$$D^E(C,D) = c(E,D) + \min_w \{D^D(C,w)\}$$

$$= 2+2 = 4$$

$$D^E(A,D) = c(E,D) + \min_w \{D^D(A,w)\}$$

$$= 2+3 = 5$$

loop!

$$D^E(A,B) = c(E,B) + \min_w \{D^B(A,w)\}$$

$$= 8+6 = 14$$

loop!

Distance table gives routing table

A	Nex t Hop	Dis t
B	-	7
C	-	∞
D	-	∞
E	-	1

B	Nex t Hop	Dis t
A	-	7
C	-	1
D	-	∞
E	-	8

C	Nex t Hop	Dis t
A	-	∞
B	-	1
D	-	2
E	-	∞

D	Nex t Hop	Dis t
A	-	∞
B	-	∞
C	-	2
E	-	2

E	Nex t Hop	Dis t
A	-	1
B	-	8
C	-	∞
D	-	2

A	Nex t Hop	Dis t
B	-	7
C	B	8
D	E	3
E	-	1

B	Nex t Hop	Dis t
A	-	7
C	-	1
D	C	3
E	-	8

C	Nex t Hop	Dis t
A	B	8
B	-	1
D	-	2
E	D	4

D	Nex t Hop	Dis t
A	E	3
B	C	3
C	-	2
E	-	2

E	Nex t Hop	Dis t
A	-	1
B	-	8
C	D	4
D	-	2

Distance table

Routing table

Distance Vector routing

A	Nex t Hop	Dis t
B	-	7
C	E	5
D	E	3
E	-	1

A	Nex t Hop	Dis t
B	E	6
C	E	5
D	E	3
E	-	1

B	Nex t Hop	Dis t
A	-	7
C	-	1
D	C	3
E	C	5

B	Nex t Hop	Dis t
A	C	6
C	-	1
D	C	3
E	C	5

C	Nex t Hop	Dis t
A	D	5
B	-	1
D	-	2
E	D	4


C	Nex t Hop	Dis t
A	D	5
B	-	1
D	-	2
E	D	4

D	Nex t Hop	Dis t
A	E	3
B	C	3
C	-	2
E	-	2


D	Nex t Hop	Dis t
A	E	3
B	C	3
C	-	2
E	-	2

E	Nex t Hop	Dis t
A	-	1
B	D	5
C	D	4
D	-	2

E	Nex t Hop	Dis t
A	-	1
B	D	5
C	D	4
D	-	2



Distance Vector



A	Nex t Hop	Dis t
B	E	6
C	E	5
D	E	3
E	-	1

B	Nex t Hop	Dis t
A	C	6
C	-	1
D	C	3
E	C	5

C	Nex t Hop	Dis t
A	D	5
B	-	1
D	-	2
E	D	4

D	Nex t Hop	Dis t
A	E	3
B	C	3
C	-	2
E	-	2

E	Nex t Hop	Dis t
A	-	1
B	D	5
C	D	4
D	-	2

Distance Vector Routing: overview

Iterative, asynchronous:

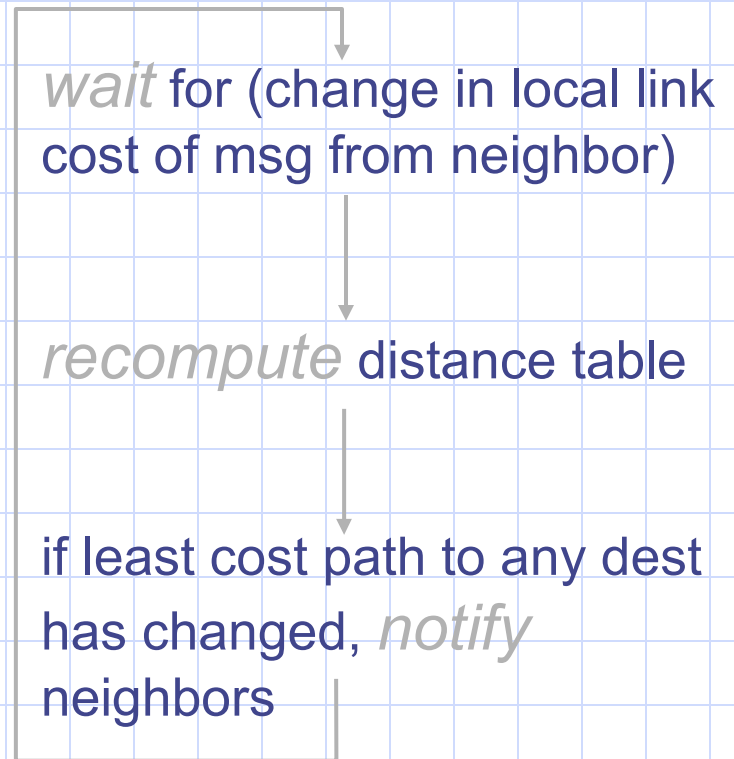
each local iteration caused by:

- ◆ local link cost change
- ◆ message from neighbor:
its least cost path change
from neighbor

Distributed:

- ◆ each node notifies
neighbors *only* when its
least cost path to any
destination changes
 - neighbors then notify their
neighbors if necessary

Each node:



Distance Vector Algorithm:

At all nodes, X:


- 1 Initialization:
- 2 for all adjacent nodes v:
- 3 $D^X(*,v) = \text{infinity}$ /* the * operator means "for all rows" */
- 4 $D^X(v,v) = c(X,v)$
- 5 for all destinations, y
- 6 send $\min_w D^X(y,w)$ to each neighbor /* w over all X's neighbors */

Distance Vector Algorithm

(cont.):

```
8 loop:
9   wait (until I see a link cost change to neighbor V
10      or until I receive update from neighbor V)
11
12   if (c(X,V) changes by d)
13     /* change cost to all dest's via neighbor v by d */
14     /* note: d could be positive or negative */
15     for all destinations y:  $D^X(y,V) = D^X(y,V) + d$ 
16
17   else if (update received from V wrt destination Y)
18     /* shortest path from V to some Y has changed */
19     /* V has sent a new value for its  $\min_w DV(Y,w)$  */
20     /* call this received new value is "newval" */
21     for the single destination y:  $D^X(Y,V) = c(X,V) + \text{newval}$ 
22
23   if we have a new  $\min_w D^X(Y,w)$  for any destination Y
24     send new value of  $\min_w D^X(Y,w)$  to all neighbors
25
26   forever
```

Distance Vector Algorithm: example



		cost via	
		Y	Z
dest	X		
	Y	2	∞
	Z	∞	7

		cost via	
		Y	Z
dest	X		
	Y	2	8
	Z	3	7

		cost via	
		Y	Z
dest	X		
	Y		
	Z		

		cost via	
		X	Z
dest	Y		
	X	2	∞
	Z	∞	1

		cost via	
		X	Z
dest	Y		
	X	2	8
	Z	9	1


		cost via	
		X	Z
dest	Y		
	X		
	Z		

		cost via	
		X	Y
dest	Z		
	X	7	∞
	Y	∞	1

		cost via	
		X	Y
dest	Z		
	X	7	3
	Y	9	1

		cost via	
		X	Y
dest	Z		
	X		
	Y		

Distance Vector Algorithm: example



		cost via	
D^X		Y	Z
d e s t	Y	2	∞
	Z	∞	7

		cost via	
D^Y		X	Z
d e s t	X	2	∞
	Z	∞	1

		cost via	
D^Z		X	Y
d e s t	X	7	∞
	Y	∞	1

		cost via	
D^X		Y	Z
d e s t	Y	2	8
	Z	3	7


$$D^X(Y,Z) = c(X,Z) + \min_w \{D^Z(Y,w)\} \\ = 7 + 1 = 8$$

$$D^X(Z,Y) = c(X,Y) + \min_w \{D^Y(Z,w)\} \\ = 2 + 1 = 3$$


Distance Vector: link cost changes

Link cost changes:

- node detects local link cost change
- updates distance table (line 15)
- if cost change in least cost path, notify neighbors (lines 23,24)



"good news travels fast"





algorithm terminates

Distance Vector: link cost changes

Link cost changes:


◆ good news travels fast

◆ bad news travels slow




Distance Vector: poisoned reverse

If Z routes through Y to get to X :



- ◆ Z tells Y its (Z's) distance to X is infinite (so Y won't route to X via Z)



Comparison of LS and DV algorithms

Message complexity

- ◆ LS: with n nodes, E links, $O(nE)$ msgs sent each
- ◆ DV: exchange between neighbors only
 - convergence time varies

Speed of Convergence

- ◆ LS: $O(n^2)$ algorithm requires $O(nE)$ msgs
 - may have oscillations
- ◆ DV: convergence time varies
 - may be routing loops
 - count-to-infinity problem

Robustness: what happens if router malfunctions?

LS:

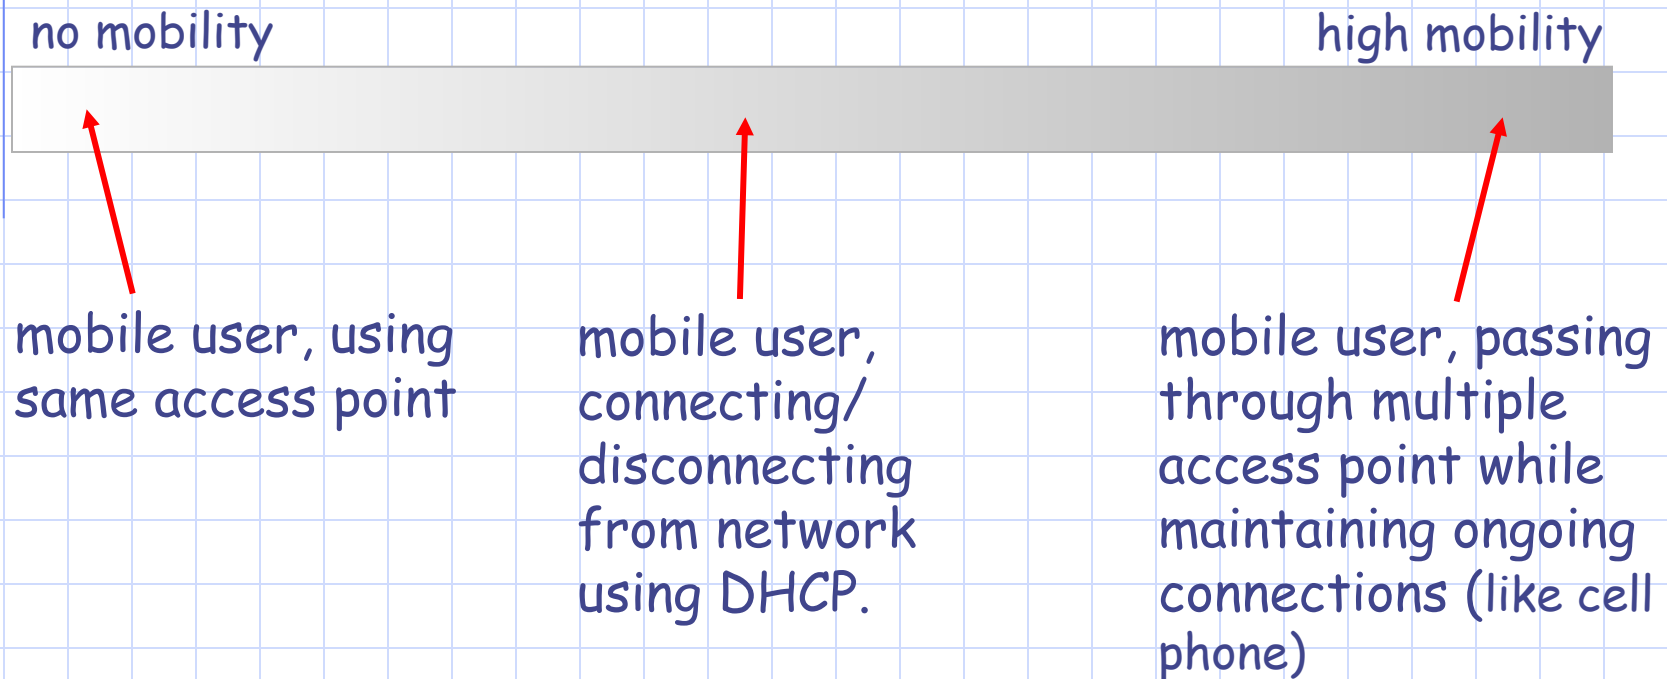
- node can advertise incorrect *link* cost
- each node computes only its *own* table

DV:

- DV node can advertise incorrect *path* cost
- each node's table used by others
 - ◆ error propagate thru network

What is mobility?

◆ spectrum of mobility, from the *network* perspective:




Mobility: Vocabulary


home network: permanent
"home" of mobile
(e.g., 128.119.40/24)

home agent: entity that will
perform mobility functions on
behalf of mobile, when mobile
is remote

Permanent address:
address in home
network, can always be
used to reach mobile
e.g., 128.119.40.186



wide area
network



correspondent

Mobility: more vocabulary

Permanent address: remains constant (e.g., 128.119.40.186)




visited network: network in which mobile currently resides (e.g., 79.129.13/24)

Care-of-address: address in visited network. (e.g., 79.129.13.2)

wide area network

correspondent: wants to communicate with mobile


home agent: entity in visited network that performs mobility functions on behalf of mobile.



How do *you* contact a mobile friend:

Consider friend frequently changing addresses, how do you find her?

- ◆ search all phone books?
- ◆ call her parents?
- ◆ expect her to let you know where he/she is?



Mobility: approaches

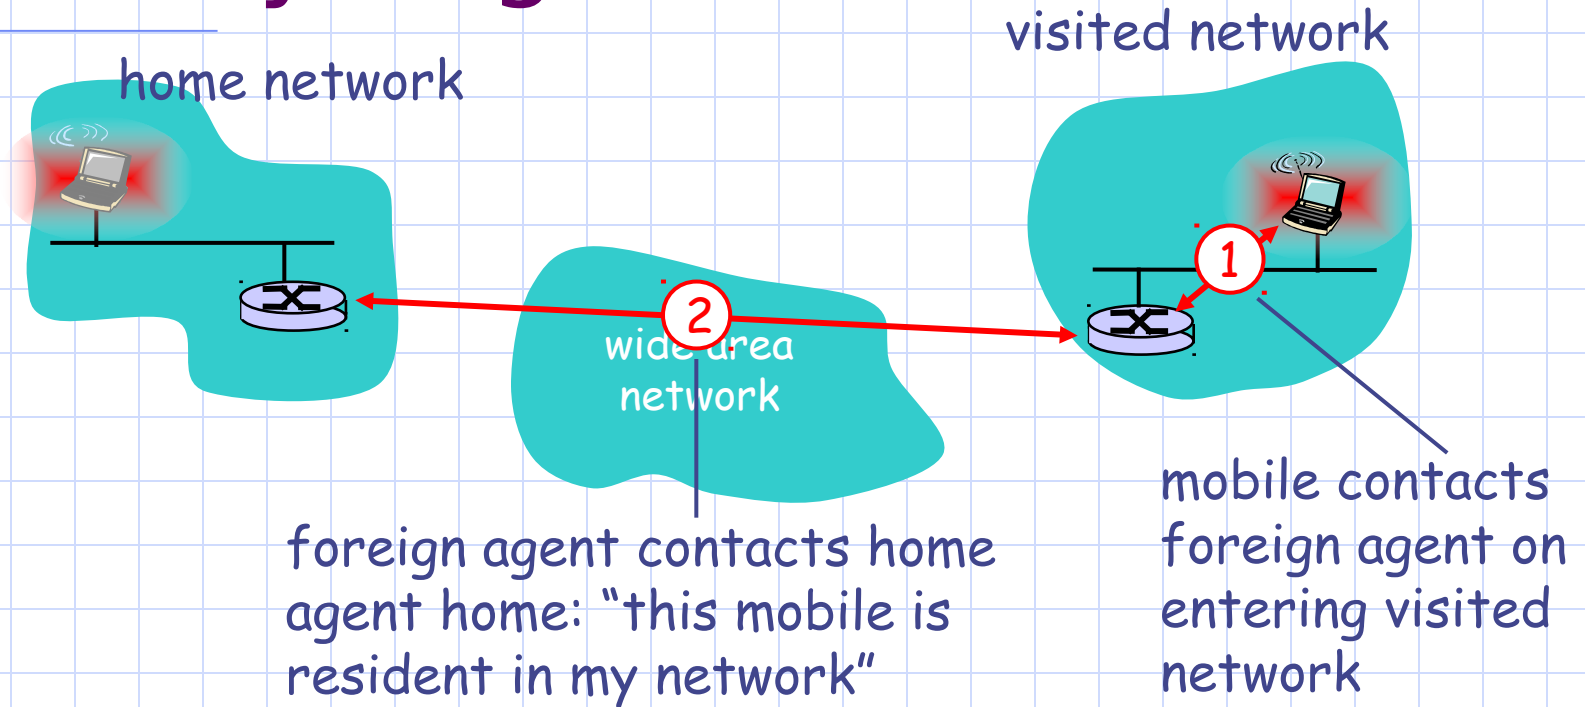
- ◆ *Let routing handle it:* routers advertise permanent address of mobile-nodes-in-residence via usual routing table exchange.
 - routing tables indicate where each mobile located
 - no changes to end-systems
- ◆ *Let end-systems handle it:*
 - *indirect routing:* communication from correspondent to mobile goes through home agent, then forwarded to remote
 - *direct routing:* correspondent gets foreign address of mobile, sends directly to mobile

Mobility: approaches

- ◆ Let routing handle it: routers advertise permanent address of mobile, mobile residence via usual routing table exchange
 - routing table entry for each mobile location
 - no changes to routing table
- ◆ *let end-systems handle it:*
 - *indirect routing:* communication from correspondent to mobile goes through home agent, then forwarded to remote
 - *direct routing:* correspondent gets foreign address of mobile, sends directly to mobile

not
scalable
to millions of
mobiles

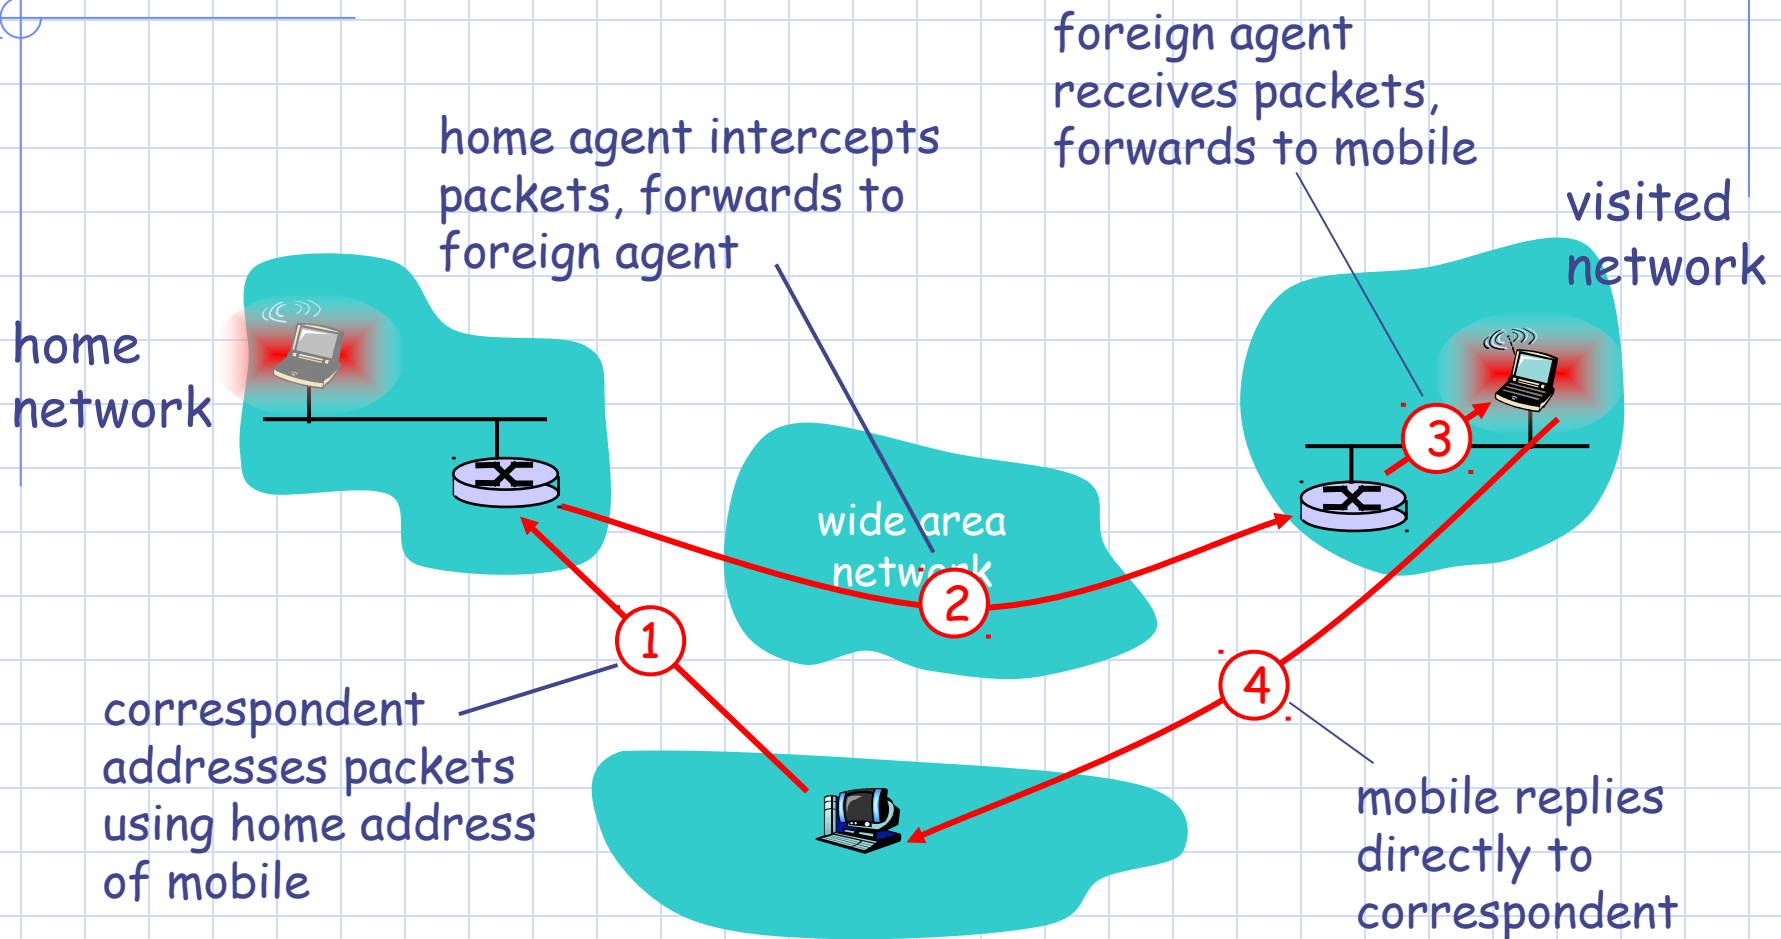
Mobility: registration



End result:


- ◆ Foreign agent knows about mobile
- ◆ Home agent knows location of mobile

Mobility via Indirect Routing

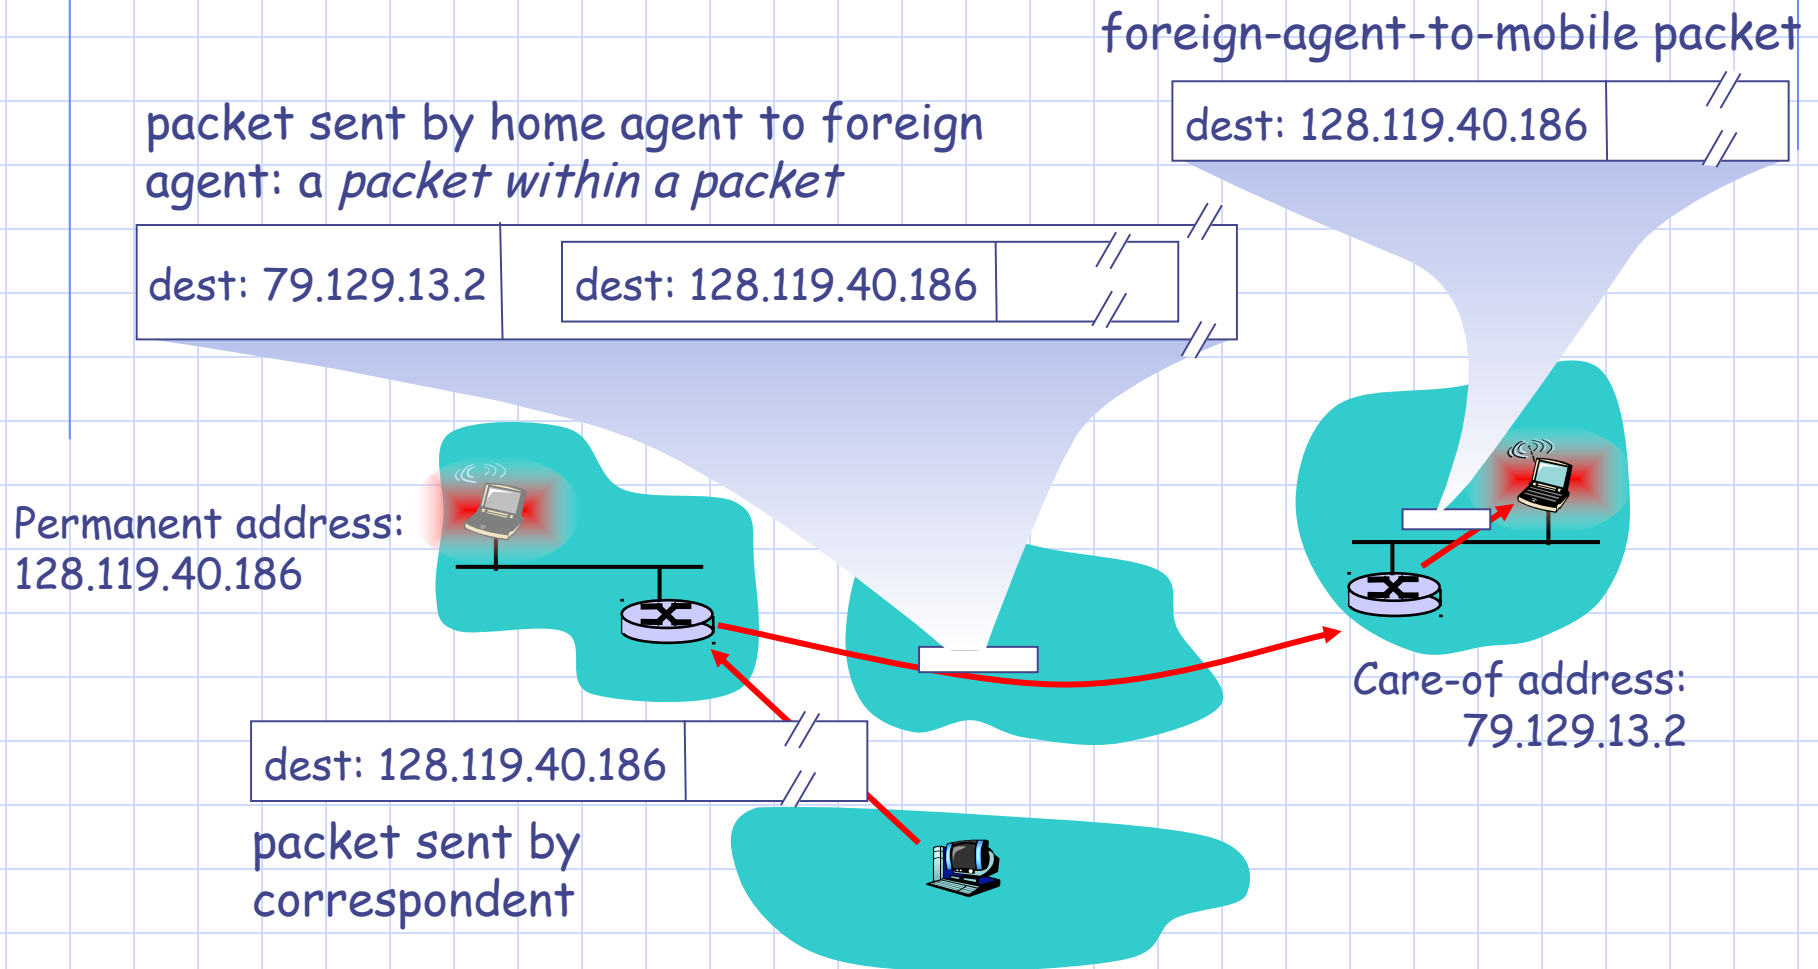


Indirect Routing: comments

- ◆ Mobile uses two addresses:
 - permanent address: used by correspondent (hence mobile location is *transparent* to correspondent)
 - care-of-address: used by home agent to forward datagrams to mobile
- ◆ foreign agent functions may be done by mobile itself
- ◆ triangle routing: correspondent-home-network-mobile
 - inefficient when correspondent, mobile are in same network



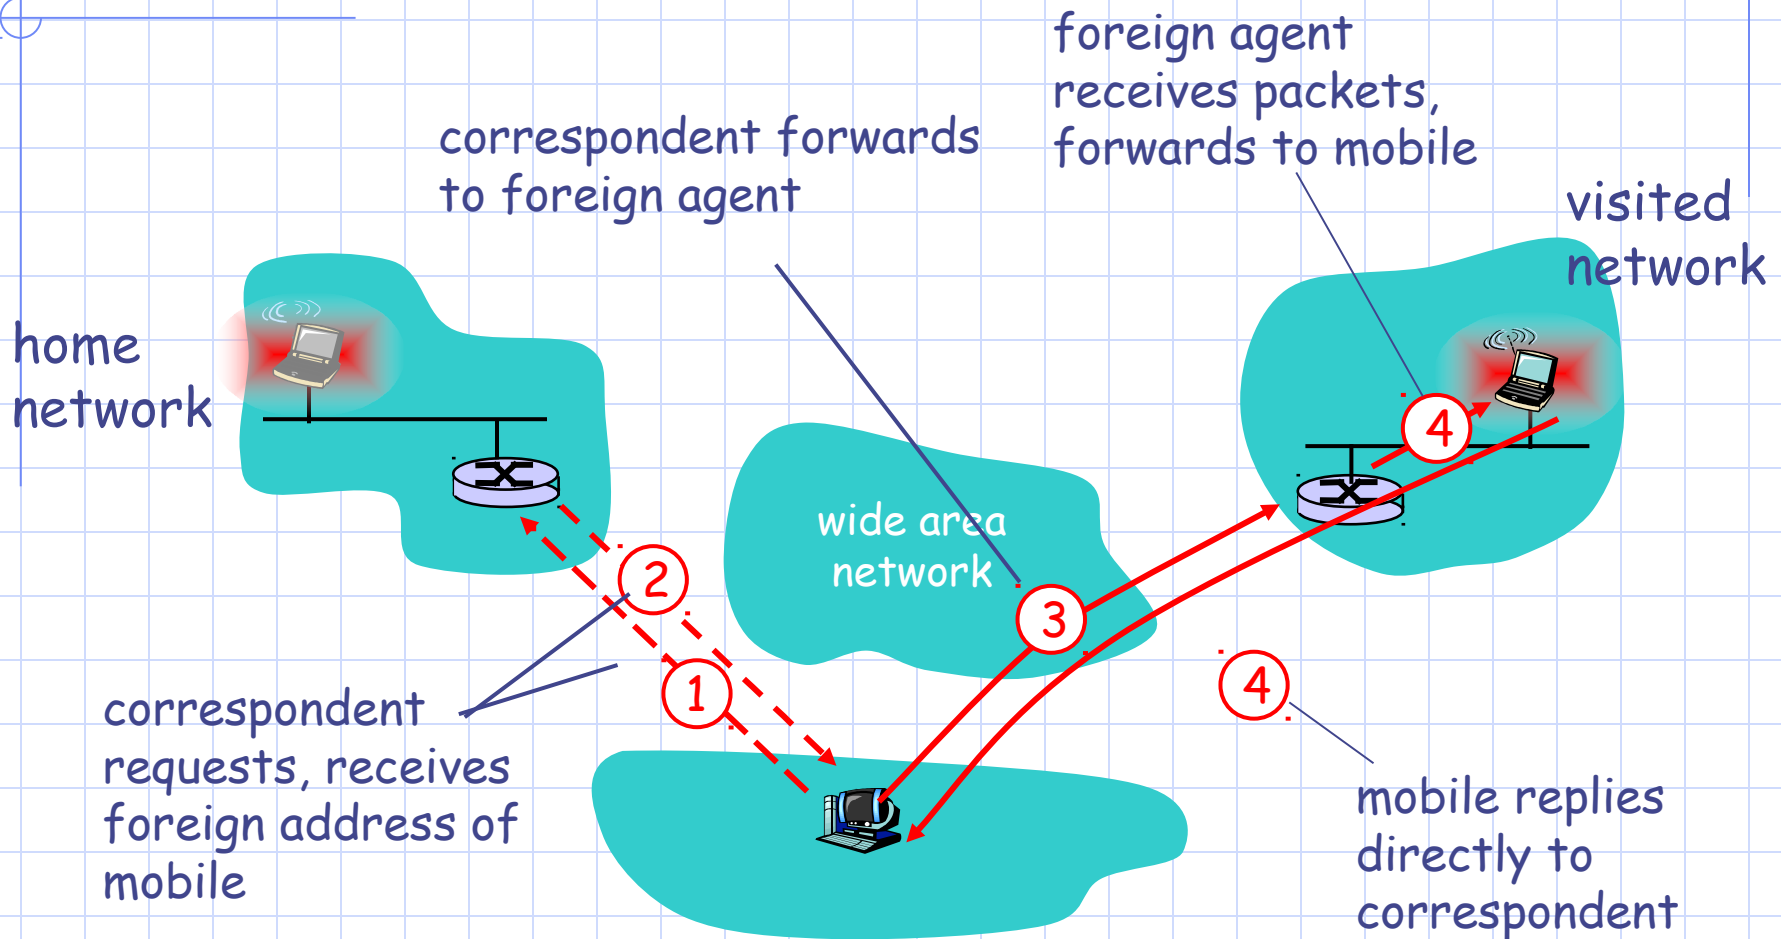
Forwarding datagrams to remote mobile



Indirect Routing: moving between networks


- ◆ suppose mobile user moves to another network
 - registers with new foreign agent
 - new foreign agent registers with home agent
 - home agent update care-of-address for mobile
 - packets continue to be forwarded to mobile (but with new care-of-address)
- ◆ Mobility, changing foreign networks transparent: *on going connections can be maintained!*

Mobility via Direct Routing



Mobility via Direct Routing: comments

- ◆ overcome triangle routing problem
- ◆ **non-transparent to correspondent:**
correspondent must get care-of-
address from home agent
 - What happens if mobile changes
network?



Mobile IP

◆ RFC 3220

◆ has many features we've seen:

- home agents, foreign agents, foreign-agent registration, care-of-addresses, encapsulation (packet-within-a-packet)

◆ three components to standard:

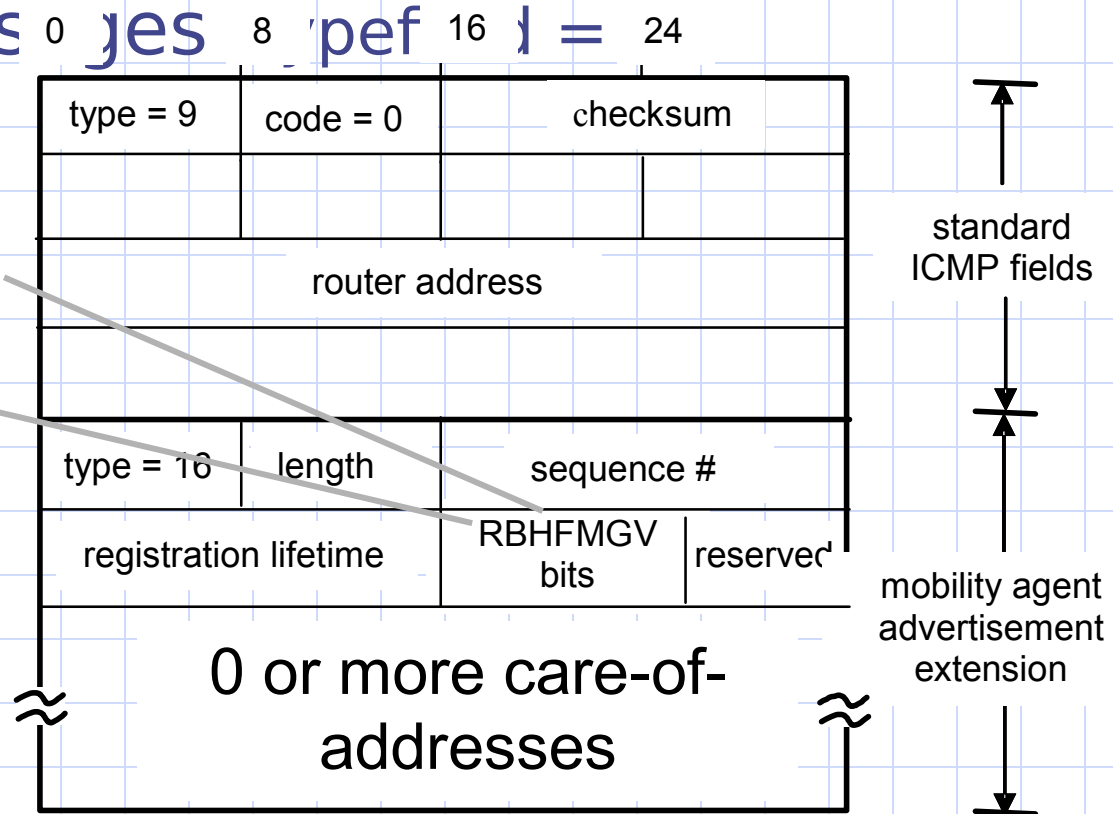
- agent discovery
- registration with home agent
- indirect routing of datagrams

Mobile IP: agent discovery

- ◆ **agent advertisement:** foreign/home agents advertise service by broadcasting ICMP mess

H,F bits: home and/or foreign agent

R bit: registration required



Mobile IP: registration example

