

Seminar 6

More on Optimization

Stored Procedures

- Advantages
 - Performance advantages
 - Server side
 - Reuse of execution plans
- Note: Requirements for plan reuse
 - Plan reuse is not always a good thing
- New in SQL Server 2005:
 - OPTIMIZE FOR / RECOMPILE query hints

Stored Procedures – Optimization Tips

- **SET NOCOUNT ON**
 - Number of affected rows is not displayed
 - Reduces network traffic
- Use **schema name** with **object name**
 - Helps finding directly the compiled plan

```
SELECT * FROM dbo.MyTable  
EXEC dbo.StoredProcedure
```

Stored Procedures – Optimization Tips

- Do not use **sp_** prefix
 - MS SQL Server first searches in **master** database and then in the **current** database
- use UNION to implement an "OR" operation
- Avoid joins between two types of columns
 - Index is not used for a converted column!

Stored Procedures – Optimization Tips

■ **sp_executesql** vs **exec**

- Execution plan of a dynamic statement can be reused if ALL characters of two consecutive executions are exactly the same

```
exec 'Select * from Categories where ID = 1'
```

```
exec 'Select * from Categories where ID = 2'
```

```
EXECUTE sp_executesql 'Select * from  
Categories where ID = @ID', '@ID int', @ID=1;
```

Stored Procedures – Optimization Tips

Cursors

- Generally use a lot of SQL Server resources and reduce the performance and scalability of your applications
- Use when you need to be able to identify scenarios where cursors are suitable/better:
 - Procedural logic / must access row by row
 - Ordered calculations

Stored Procedures – Optimization Tips

- Do not use COUNT() in a subquery to do an existence check
- Use IF EXISTS (SELECT 1 FROM...)
 - The output of nested select is not used
 - Reduces processing time and network transfer
- Keep transaction short
 - Transaction length affects blocking and deadlocking

Stored Procedures – Optimization Tips

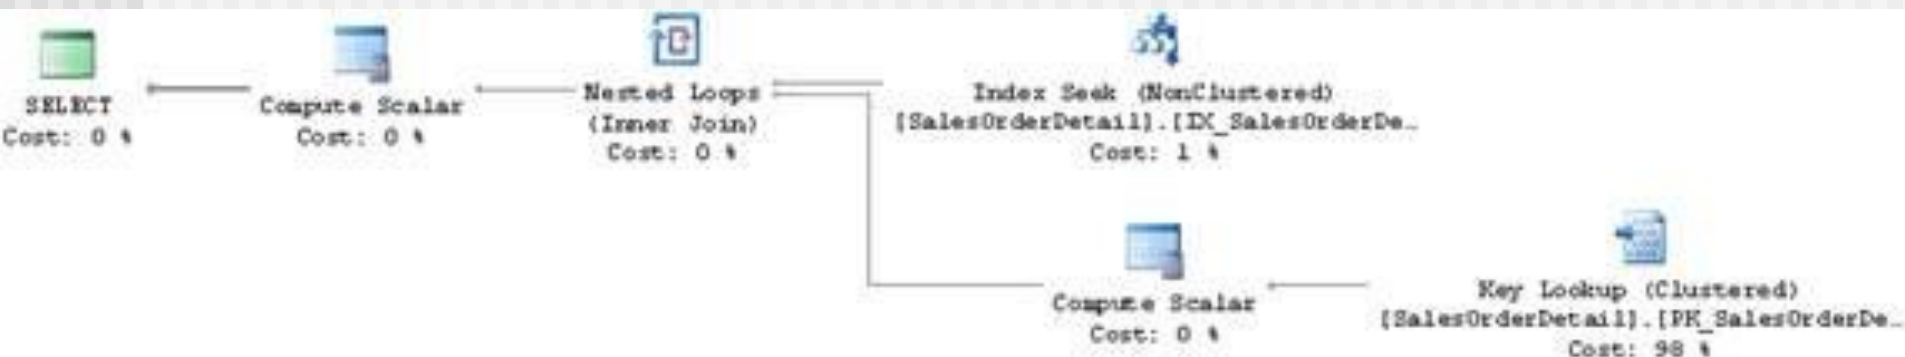
Reuse execution plan

```
CREATE PROCEDURE test (@pid int)
```

```
AS
```

```
    SELECT * FROM Sales.SalesOrderDetail  
    WHERE ProductID = @pid
```

```
exec test(897)
```



```
exec test(870)
```



Stored Procedures – Optimization Tips

OPTIMIZE FOR / RECOMPILE query hints

```
ALTER PROCEDURE test (@pid int)
AS
    SELECT * FROM Sales.SalesOrderDetail
    WHERE ProductID = @pid
    OPTION (OPTIMIZE FOR (@pid = 870))
```

```
ALTER PROCEDURE test (@pid int)
AS
    SELECT * FROM Sales.SalesOrderDetail
    WHERE ProductID = @pid
    OPTION (RECOMPILE)
```

Stored Procedures – Optimization Tips

OPTIMIZE FOR UNKNOWN

- local variables are not known at optimization time
- example below: always generates the same execution plan

```
ALTER PROCEDURE test (@pid int)
AS
    DECLARE @lpid int
    SET @lpid = @pid
    SELECT * FROM Sales.SalesOrderDetail
    WHERE ProductID = @lpid
```

Stored Procedures – Optimization Tips

OPTIMIZE FOR UNKNOWN

- local variables are not known at optimization time
- example below: always generates the same execution plan

```
ALTER PROCEDURE test (@pid int)
AS
    SELECT * FROM Sales.SalesOrderDetail
    WHERE ProductID = @pid
    OPTION (OPTIMIZE FOR UNKNOWN)
```

Stored Procedures – Optimization Tips

OPTIMIZE FOR query hints

```
DECLARE @city_name nvarchar(30);  
DECLARE @postal_code nvarchar(15);  
  
SELECT * FROM Person.Address  
WHERE City = @city_name AND PostalCode =  
@postal_code OPTION  
(OPTIMIZE FOR (@city_name = 'Seattle',  
@postal_code UNKNOWN) );
```

Stored Procedures – Optimization Tips

Other query hints

- HASH GROUP vs ORDER GROUP

```
SELECT ProductID, OrderQty, SUM(LineTotal) AS Total
FROM Sales.SalesOrderDetail
WHERE UnitPrice < $5.00
GROUP BY ProductID, OrderQty
ORDER BY ProductID, OrderQty
OPTION (HASH GROUP, FAST 10);
```

Stored Procedures – Optimization Tips

Other query hints

- MERGE UNION vs
HASH UNION vs
CONCAT UNION

```
SELECT ...  
UNION  
SELECT ...
```

```
OPTION ( MERGE UNION )
```

Stored Procedures – Optimization Tips

Join hints

- LOOP JOIN vs
MERGE JOIN vs
HASH JOIN

```
SELECT * FROM Sales.Customer AS c
INNER JOIN Sales.vStoreWithAddresses AS sa
        ON c.CustomerID = sa.BusinessEntityID
WHERE TerritoryID = 5
OPTION (MERGE JOIN);
GO
```

Stored Procedures – Optimization Tips

Join hints

- FAST n - focus on returning the first 'n' rows as fast as possible

```
SELECT * FROM Sales.Customer AS c
INNER JOIN Sales.vStoreWithAddresses AS sa
        ON c.CustomerID = sa.BusinessEntityID
WHERE TerritoryID = 5
OPTION (FAST 10);
GO
```


Stored Procedures – Optimization Tips

Join hints

- FORCE ORDER – “force” the optimizer to use the order of joins as they are listed in the query

```
SELECT * FROM Table1  
INNER JOIN Table2 ON Table1.a = Table2.b  
INNER JOIN Table3 ON Table2.c = Table3.d  
INNER JOIN Table4 ON Table3.e = Table4.f  
OPTION (FORCE ORDER);
```

Stored Procedures – Optimization Tips

More about

Controlling Execution Plans with Hints

<https://www.simple-talk.com/sql/performance/controlling-execution-plans-with-hints/>

Dynamic Execution

- Disadvantages:
 - Ugly code; hard to maintain
 - Requires direct permissions (in 2000)
 - Security risk of SQL Injection
- Use smartly:
 - Dynamic filters and sorting to get good plans
 - And more....

Temporary Tables

- Useful when:
 - You have intermediate result sets that need to be accessed several times
 - You need a temporary storage area for data while running procedural code
- Use temp tables when:
 - Typically large volumes of data, where plan efficiency is important and non-trivial
- Use table variables when:
 - Typically small volumes of data, where plan efficiency is not as significant as recompilations, or when plans are trivial

Triggers

- Expensive (when rollback, undo as opposed to reject)
- Main performance impact involves accessing inserted and deleted views
 - SQL Server 2000: transaction log
 - SQL Server 2005: row versioning (tempdb)
- Stripe transaction log/tempdb when using triggers
- Try to utilize set-based activities
- Identify # of affected rows and react accordingly
- UPDATE triggers record delete followed by insert in the log

SQL Server Options

