



Seminar 6

Indexes Best practices (II)

General Index Design Guidelines

■ Database considerations

- Too many indexes on a table affect the performance of INSERT, UPDATE, DELETE, MERGE statements
- Indexing small tables may not be optimal
- Indexes on views are useful when views contain aggregations and/or table joins

■ Query considerations

- Create nonclustered indexes for columns frequently used in WHEREs and JOINS
- Covering indexes can improve query performance
- Write queries that insert or modify as many rows as possible in a single statement
- Evaluate the query type and how columns are used in the query

General Index Design Guidelines

- Column Considerations
 - Keep the length of index key short for clustered indexes
 - Clustered indexes are better on unique/nonnull cols
 - Columns of **ntext**, **text**, **image**, **varchar(max)**, **nvarchar(max)**, **varbinary(max)** cannot be specified as index key columns
 - Examine column uniqueness
 - Examine data distribution in column (avoid indexes on columns with few unique values) – use filtered indexes
 - Consider the order of the columns for multiple index. Columns used in an equal to (=), greater than (>), less than (<), or BETWEEN search condition should be placed first. Additional columns should be ordered from the most distinct to the least distinct.
 - Consider indexing computed columns.

Unique Indexes

- A unique index guarantees that the index key contains no duplicate values
- Specifying a unique index makes sense only when key columns are unique
- Uniqueness – helpful information for query optimizer

Filtered Indexes

Filtered Index: an optimized nonclustered index, especially suited to cover queries that select from a well-defined subset of data

```
CREATE NONCLUSTERED INDEX FI_EndDate ON  
Products (ProductID, EndDate)  
WHERE EndDate IS NOT NULL ;  
GO
```

- Improved query performance
- Reduced index maintenance costs
- Reduced index storage costs

Indexes for Deletes

At DELETE:

- SQL Server will check for dependent rows by examining all foreign keys
- It will then check any related tables for data.
 - If there is an index, SQL Server will use that index to check for related data
 - If there isn't an index, though, SQL Server will have to **scan** the table for data.
- Deletes could be very slow if there is no index defined for foreign keys

Indexed Views

SET options	Required value	Default server value
ANSI_NULLS	ON	ON
ANSI_PADDING	ON	ON
ANSI_WARNINGS	ON	ON
ARITHABORT	ON	ON
CONCAT_NULL_YIELDS_NULL	ON	ON
NUMERIC_ROUNDABORT	OFF	OFF
QUOTED_IDENTIFIER	ON	ON

Indexed Views Restrictions

- SELECT statement cannot reference other views
- All functions must be deterministic
- AVG, MIN, MAX, STDEV, STDEVP, VAR and VARP are not allowed
- The index must be both clustered and unique
- SELECT statement must not contain subqueries, outer joins, EXCEPT, INTERSECT, TOP, UNION, ORDER BY, DISTINCT etc

Columnstore Indexes

- groups and stores data for each column and then joins all the columns to complete the whole index
- Suited for warehouses (read only tables)

Hard and fast rules for indexing

- Each table should have a clustered index that is (ideally) small, selective, ever increasing, and static. (a table without a clustered index is called a *heap*.)
- Implement nonclustered indexes on foreign key relationships
- Implement nonclustered indexes on columns that are frequently used in WHERE clauses.
- Do not implement single-column indexes on every column in a table. This will cause high overhead.
- In multi-column indexes, list the most selective (nearest to unique) first in the column list.
- For most often-used queries create covering nonclustered index.

Fragmentation

- *Internal Fragmentation*: records are stored non-contiguously inside the page. Internal fragmentation occurs if there is unused space between records in a page. The fullness of each page can vary over time. This unused space causes poor cache utilization and more I/O, which ultimately leads to poor query performance.
- *External Fragmentation*: When on disk, the physical storage of pages and extents is not contiguous. When the extents of a table are not physically stored contiguously on disk, switching from one extent to another causes higher disk rotations.

Fragmentation

- *Logical Fragmentation*: Every index page is linked with previous and next page in the logical order of column data. Because of Page Split, the pages turn into *out-of-order* pages.
- An *out-of-order* page is a page for which the next physical page allocated to the index is not the page pointed to by the next-page pointer in the current leaf page.


Page read requests: 2


Extent switches: 0

Disk space used by table: 16 KB

avg_fragmentation_in_percent: 0

avg_page_space_used_in_percent: 100





Page read requests: 6

Extent switches: 5

Disk space used by table:
48 KB

avg_fragmentation_in_
percent > 80

avg_page_space_used_in_
percent: 33

Fragmentation

- *sys.dm_db_index_physical_stats*
 - **avg_fragmentation_in_percent:** This is a percentage value that represents external fragmentation.
 - **avg_page_space_used_in_percent:** This is an average percentage use of pages that represents to internal fragmentation.
- **Reducing Fragmentation in a Heap:**
 - To reduce the fragmentation of a heap, create a clustered index on the table.
 - Creating the clustered index: rearrange the records in an order, and then place the pages contiguously on disk.

Fragmentation

Reducing Fragmentation in a Index:

- If `avg_fragmentation_in_percent` $> 5\%$ and $< 30\%$, then use `ALTER INDEX REORGANIZE`:
 - reorder the leaf level pages of the index in a logical order.
- If `avg_fragmentation_in_percent` $> 30\%$, then use `ALTER INDEX REBUILD`:
 - replacement for `DBCC DBREINDEX` to rebuild the index online or offline. In such case, we can also use the drop and re-create index method.
- Drop and re-create the clustered index:
 - Re-creating a clustered index redistributes the data and results in full data pages. The level of fullness can be configured by using the `FILLFACTOR` option in `CREATE INDEX`.