



Buy in [print](#) and [eBook](#).

Table of Contents

Prologue

I. Language Concepts

1. A Guided Tour
2. Variables and Functions
3. Lists and Patterns
4. Files, Modules, and Programs
5. Records
6. Variants
7. Error Handling
- 8. Imperative Programming**
9. Functors
10. First-Class Modules
11. Objects
12. Classes

II. Tools and Techniques

III. The Runtime System

Index

Chapter 8. Imperative Programming

Most of the code shown so far in this book, and indeed, **most OCaml code in general, is *pure*. Pure code works without mutating the program's internal state, performing I/O, reading the clock, or in any other way interacting with changeable parts of the world.** Thus, **a pure function behaves like a mathematical function, always returning the same results when given the same inputs, and never affecting the world except insofar as it returns the value of its computation.** Imperative code, on the other hand, operates by side effects that modify a program's internal state or interact with the outside world. An imperative function has a new effect, and potentially returns different results, every time it's called.[O comments](#)

Pure code is the default in OCaml, and for good reason—it's generally easier to reason about, less error prone and more composable. But imperative code is of fundamental importance to any practical programming language, because real-world tasks require that you interact with the outside world, which is by its nature imperative. Imperative programming can also be important for performance. While pure code is quite efficient in OCaml, there are many algorithms that can only be implemented efficiently using imperative techniques.[O comments](#)

OCaml offers a happy compromise here, making it easy and natural to program in a pure style, but also providing great support for imperative programming. This chapter will walk you through OCaml's imperative features, and help you use them to their fullest.[O comments](#)

EXAMPLE: IMPERATIVE DICTIONARIES

We'll start with the implementation of a simple imperative dictionary, i.e., a mutable mapping from keys to values. This is really for illustration purposes; both Core and the standard library provide imperative dictionaries, and for most real-world tasks, you should use one of those implementations. There's more advice on using Core's implementation in particular in [Chapter 13, Maps and Hash Tables](#).[O comments](#)

The dictionary we'll describe now, like those in Core and the standard library, will be implemented as a hash table. In particular, we'll use an *open hashing* scheme, where the hash table will be an array of buckets, each bucket containing a list of key/value pairs that have been hashed into that bucket.[O comments](#)

Here's the interface we'll match, provided as an `mli`. The type `('a, 'b) t` represents a dictionary with keys of type `'a` and data of type `'b`:[O comments](#)

```
(* file: dictionary.mli *)
open Core.Std

type ('a, 'b) t

val create : unit -> ('a, 'b) t
val length : ('a, 'b) t -> int
val add : ('a, 'b) t -> key:'a -> data:'b -> unit
val find : ('a, 'b) t -> 'a -> 'b option
val iter : ('a, 'b) t -> f:(key:'a -> data:'b -> unit) -> unit
val remove : ('a, 'b) t -> 'a -> unit
```

OCaml * imperative-programming/dictionary.mli , continued (part 1) * all code

The `mli` also includes a collection of helper functions whose purpose and behavior should be largely inferable from their names and type signatures. Notice that a number of the functions, in particular, ones like `add` that modify the dictionary, return `unit`. This is typical of functions that act by side effect.[O comments](#)

We'll now walk through the implementation (contained in the corresponding `ml` file) piece by piece, explaining different imperative constructs as they come up.[O comments](#)

Our first step is to define the type of a dictionary as a record with two fields:[O comments](#)

```
(* file: dictionary.ml *)
open Core.Std

type ('a, 'b) t = { mutable length: int;
                   buckets: ('a * 'b) list array;
                   }

OCaml * imperative-programming/dictionary.ml , continued (part 1) * all code
```

The first field, `length`, is declared as mutable. In OCaml, records are immutable by default, but individual fields are mutable when marked as such. The second field, `buckets`, is immutable but contains an array, which is itself a mutable data structure.[0 comments](#)

Now we'll start putting together the basic functions for manipulating a dictionary:[0 comments](#)

```
let num_buckets = 17

let hash_bucket key = (Hashtbl.hash key) mod num_buckets

let create () =
{ length = 0;
  buckets = Array.create ~len:num_buckets [];
}

let length t = t.length

let find t key =
  List.find_map t.buckets.(hash_bucket key)
  ~f:(fun (key', data) -> if key' = key then Some data else None)
```

OCaml * imperative-programming/dictionary.ml , continued (part 2) * all code

Note that `num_buckets` is a constant, which means our bucket array is of fixed length. A practical implementation would need to be able to grow the array as the number of elements in the dictionary increases, but we'll omit this to simplify the presentation.[0 comments](#)

The function `hash_bucket` is used throughout the rest of the module to choose the position in the array that a given key should be stored at. It is implemented on top of `Hashtbl.hash`, which is a hash function provided by the OCaml runtime that can be applied to values of any type. Thus, its own type is polymorphic: `'a -> int`.[0 comments](#)

The other functions defined above are fairly straightforward:[0 comments](#)

`create`

Creates an empty dictionary.[0 comments](#)

`length`

Grabs the length from the corresponding record field, thus returning the number of entries stored in the dictionary.[0 comments](#)

`find`

Looks for a matching key in the table and returns the corresponding value if found as an option.[0 comments](#)

Another important piece of imperative syntax shows up in `find`: we write `array.(index)` to grab a value from an array. `find` also uses `List.find_map`, which you can see the type of by typing it into the toplevel:[0 comments](#)

```
# List.find_map;;
- : 'a list -> f:('a -> 'b option) -> 'b option = <fun>
```

OCaml Utop * imperative-programming/examples.topscript , continued (part 1) * all code

`List.find_map` iterates over the elements of the list, calling `f` on each one until a `Some` is returned by `f`, at which point that value is returned. If `f` returns `None` on all values, then `None` is returned.[0 comments](#)

Now let's look at the implementation of `iter`:[0 comments](#)

```
let iter t ~f =
  for i = 0 to Array.length t.buckets - 1 do
    List.iter t.buckets.(i) ~f:(fun (key, data) -> f ~key ~data)
  done
```

OCaml * imperative-programming/dictionary.ml , continued (part 3) * all code

`iter` is designed to walk over all the entries in the dictionary. In particular, `iter t ~f` will call `f` for each key/value pair in dictionary `t`. Note that `f` must return `unit`, since it is expected to work by side effect rather than by returning a value, and the overall `iter` function returns `unit` as well.[0 comments](#)

The code for `iter` uses two forms of iteration: a `for` loop to walk over the array of buckets; and within that loop a call to `List.iter` to walk over the values in a given bucket. We could have done the outer loop with a recursive function instead of a `for` loop, but `for` loops are syntactically convenient, and are more familiar and idiomatic in imperative contexts.[0 comments](#)

The following code is for adding and removing mappings from the dictionary:[0 comments](#)

```
let bucket_has_key t i key =
  List.exists t.buckets.(i) ~f:(fun (key',_) -> key' = key)

let add t ~key ~data =
  let i = hash_bucket key in
  let replace = bucket_has_key t i key in
  let filtered_bucket =
    if replace then
      List.filter t.buckets.(i) ~f:(fun (key',_) -> key' <> key)
    else
      t.buckets.(i)
  in
  t.buckets.(i) <- (key, data) :: filtered_bucket;
  if not replace then t.length <- t.length + 1

let remove t key =
  let i = hash_bucket key in
  if bucket_has_key t i key then (
    let filtered_bucket =
      List.filter t.buckets.(i) ~f:(fun (key',_) -> key' <> key)
    in
    t.buckets.(i) <- filtered_bucket;
    t.length <- t.length - 1
  )
)
```

OCaml * imperative-programming/dictionary.ml , continued (part 4) * all code

This preceding code is made more complicated by the fact that we need to detect whether we are overwriting or removing an existing binding, so we can decide whether `t.length` needs to be changed. The helper function `bucket_has_key` is used for this purpose.[0 comments](#)

Another piece of syntax shows up in both `add` and `remove`: the use of the `<-` operator to update elements of an array (`array.(i) <- expr`) and for updating a record field (`record.field <- expression`).[0 comments](#)

We also use `;`, the sequencing operator, to express a sequence of imperative actions. We could have done the same using `let` bindings:[0 comments](#)

```
let () = t.buckets.(i) <- (key, data) :: filtered_bucket in
  if not replace then t.length <- t.length + 1
```

OCaml * imperative-programming/dictionary2.ml , continued (part 1) * all code

but `;` is more concise and idiomatic. More generally,[0 comments](#)

```
<expr1>;
<expr2>;
...
<exprN>
```

Syntax * imperative-programming/semicolon.syntax * all code

is equivalent to[0 comments](#)

```
let () = <expr1> in
let () = <expr2> in
...
<exprN>
```

Syntax * imperative-programming/let-unit.syntax * all code

When a sequence expression `expr1; expr2` is evaluated, `expr1` is evaluated first, and then `expr2`. The expression `expr1` should have type `unit` (though this is a warning rather than a hard restriction. The `-strict-sequence` compiler flag makes this a hard restriction, which is generally a good idea), and the value of `expr2` is returned as the value of the entire sequence. For example, the sequence `print_string "hello world"; 1 + 2` first prints the string "hello world", then returns the integer 3.[0 comments](#)

Note also that we do all of the side-effecting operations at the very end of each function. This is good practice because it minimizes the chance that such operations will be interrupted with an exception, leaving the data structure in an inconsistent state.[0 comments](#)

PRIMITIVE MUTABLE DATA

Now that we've looked at a complete example, let's take a more systematic look at imperative

programming in OCaml. We encountered two different forms of mutable data above: records with mutable fields and arrays. We'll now discuss these in more detail, along with the other primitive forms of mutable data that are available in OCaml.[0 comments](#)

Array-Like Data

OCaml supports a number of array-like data structures; i.e., mutable integer-indexed containers that provide constant-time access to their elements. We'll discuss several of them in this section.[0 comments](#)

Ordinary arrays

The `array` type is used for general-purpose polymorphic arrays. The `Array` module has a variety of utility functions for interacting with arrays, including a number of mutating operations. These include `Array.set`, for setting an individual element, and `Array.blit`, for efficiently copying values from one range of indices to another.[0 comments](#)

Arrays also come with special syntax for retrieving an element from an array:[0 comments](#)

```
<array_expr>.[<index_expr>]
```

Syntax * imperative-programming/array-get.syntax * all code

and for setting an element in an array:[0 comments](#)

```
<array_expr>.[<index_expr>] <- <value_expr>
```

Syntax * imperative-programming/array-set.syntax * all code

Out-of-bounds accesses for arrays (and indeed for all the array-like data structures) will lead to an exception being thrown.[0 comments](#)

Array literals are written using `[]` and `[]` as delimiters. Thus, `[1; 2; 3]` is a literal integer array.[0 comments](#)

Strings

Strings are essentially byte arrays which are often used for textual data. The main advantage of using a `string` in place of a `Char.t` array (a `Char.t` is an 8-bit character) is that the former is considerably more space-efficient; an array uses one word—8 bytes on a 64-bit machine—to store a single entry, whereas strings use 1 byte per character.[0 comments](#)

Strings also come with their own syntax for getting and setting values:[0 comments](#)

```
<string_expr>.[<index_expr>]  
<string_expr>.[<index_expr>] <- <char_expr>
```

Syntax * imperative-programming/string.syntax * all code

And string literals are bounded by quotes. There's also a module `String` where you'll find useful functions for working with strings.[0 comments](#)

Bigarrays

A `Bigarray.t` is a handle to a block of memory stored outside of the OCaml heap. These are mostly useful for interacting with C or Fortran libraries, and are discussed in [Chapter 20](#), *Memory Representation of Values*. Bigarrays too have their own getting and setting syntax:[0 comments](#)

```
<bigarray_expr>.{<index_expr>}  
<bigarray_expr>.{<index_expr>} <- <value_expr>
```

Syntax * imperative-programming/bigarray.syntax * all code

Mutable Record and Object Fields and Ref Cells

As we've seen, records are immutable by default, but individual record fields can be declared as mutable. These mutable fields can be set using the `<-` operator, i.e., `record.field <- expr`.[0 comments](#)

As we'll see in [Chapter 11](#), *Objects*, fields of an object can similarly be declared as mutable, and can then be modified in much the same way as record fields.[0 comments](#)

Ref cells

Variables in OCaml are never mutable—they can refer to mutable data, but what the variable points to can't be changed. Sometimes, though, you want to do exactly what you would do with a mutable variable in another language: define a single, mutable value. In OCaml this is typically achieved using a `ref`, which is essentially a container with a single mutable polymorphic field.[0 comments](#)

The definition for the `ref` type is as follows:[0 comments](#)

```
# type 'a ref = { mutable contents : 'a };;  
type 'a ref = { mutable contents : 'a; }  
  
OCaml Utop * imperative-programming/ref.topscript , continued (part 1) * all code
```

The standard library defines the following operators for working with `refs`.[0 comments](#)

`ref expr`
Constructs a reference cell containing the value defined by the expression `expr`.[0 comments](#)

`!refcell`
Returns the contents of the reference `cell`.[0 comments](#)

`refcell := expr`
Replaces the contents of the reference `cell`.[0 comments](#)

You can see these in action:[0 comments](#)

```
# let x = ref 1;;  
val x : int ref = {contents = 1}  
# !x;;  
- : int = 1  
# x := !x + 1;;  
- : unit = ()  
# !x;;  
- : int = 2  
  
OCaml Utop * imperative-programming/ref.topscript , continued (part 3) * all code
```

The preceding are just ordinary OCaml functions, which could be defined as follows:[0 comments](#)

```
# let ref x = { contents = x };;  
val ref : 'a -> 'a ref = <fun>  
# let (!) r = r.contents;;  
val ( ! ) : 'a ref -> 'a = <fun>  
# let (:=) r x = r.contents <- x;;  
val ( := ) : 'a ref -> 'a -> unit = <fun>  
  
OCaml Utop * imperative-programming/ref.topscript , continued (part 2) * all code
```

Foreign Functions

Another source of imperative operations in OCaml is resources that come from interfacing with external libraries through OCaml's foreign function interface (FFI). The FFI opens OCaml up to imperative constructs that are exported by system calls or other external libraries. Many of these come built in, like access to the `write` system call or to the `clock`, while others come from user libraries, like LAPACK bindings. OCaml's FFI is discussed in more detail in [Chapter 19, Foreign Function Interface](#).[0 comments](#)

FOR AND WHILE LOOPS

OCaml provides support for traditional imperative looping constructs, in particular, `for` and `while` loops. Neither of these constructs is strictly necessary, since they can be simulated with recursive functions. Nonetheless, explicit `for` and `while` loops are both more concise and more idiomatic when programming imperatively.[0 comments](#)

The `for` loop is the simpler of the two. Indeed, we've already seen the `for` loop in action—the `iter` function in `Dictionary` is built using it. Here's a simple example of `for`:[0 comments](#)

```
# for i = 0 to 3 do printf "i = %d\n" i done;;  
  
i = 0  
i = 1
```

```
i = 2
i = 3
- : unit = ()
```

OCaml Utop * imperative-programming/for.topscript , continued (part 1) * all code

As you can see, the upper and lower bounds are inclusive. We can also use `downto` to iterate in the other direction:[0 comments](#)

```
# for i = 3 downto 0 do printf "i = %d\n" i done;;

i = 3
i = 2
i = 1
i = 0
- : unit = ()
```

OCaml Utop * imperative-programming/for.topscript , continued (part 2) * all code

Note that the loop variable of a `for` loop, `i` in this case, is immutable in the scope of the loop and is also local to the loop, i.e., it can't be referenced outside of the loop.[0 comments](#)

OCaml also supports `while` loops, which include a condition and a body. The loop first evaluates the condition, and then, if it evaluates to true, evaluates the body and starts the loop again. Here's a simple example of a function for reversing an array in place:[0 comments](#)

```
# let rev_inplace ar =
  let i = ref 0 in
  let j = ref (Array.length ar - 1) in
  (* terminate when the upper and lower indices meet *)
  while !i < !j do
    (* swap the two elements *)
    let tmp = ar.(!i) in
    ar.(!i) <- ar.(!j);
    ar.(!j) <- tmp;
    (* bump the indices *)
    incr i;
    decr j
  done
;;
val rev_inplace : 'a array -> unit = <fun>
# let nums = [|1;2;3;4;5|];;
val nums : int array = [|1; 2; 3; 4; 5|]
# rev_inplace nums;;
- : unit = ()
# nums;;
- : int array = [|5; 4; 3; 2; 1|]
```

OCaml Utop * imperative-programming/for.topscript , continued (part 3) * all code

In the preceding example, we used `incr` and `decr`, which are built-in functions for incrementing and decrementing an `int ref` by one, respectively.[0 comments](#)

EXAMPLE: DOUBLY LINKED LISTS

Another common imperative data structure is the doubly linked list. Doubly linked lists can be traversed in both directions, and elements can be added and removed from the list in constant time. Core defines a doubly linked list (the module is called `Doubly_linked`), but we'll define our own linked list library as an illustration.[0 comments](#)

Here's the `mli` of the module we'll build:[0 comments](#)

```
(* file: dlist.mli *)
open Core.Std

type 'a t
type 'a element

(** Basic List operations *)
val create : unit -> 'a t
val is_empty : 'a t -> bool

(** Navigation using [element]s *)
val first : 'a t -> 'a element option
val next : 'a element -> 'a element option
val prev : 'a element -> 'a element option
val value : 'a element -> 'a
```

OCaml * imperative-programming/dlist.mli * all code

Now let's look at the implementation. We'll start by defining 'a element and 'a t:0 comments

OCaml * imperative-programming/dlist.ml , continued (part 1) * all code

Now we can define a few basic functions that operate on lists and elements:[0 comments](#)

OCaml * imperative-programming/dlist.ml , continued (part 2) * all code

```
# let rec endless_loop = 1 :: 2 :: 3 :: endless_loop;;  
val endless_loop : int list =  
[1; 2; 3; 1; 2; 3; 1; 2; 3; 1; 2; 3; 1; 2; 3; 1; 2; 3; 1; 2; 3;  
 2; 3; 1; 2; 3; 1; 2; 3; 1; 2; 3; 1; 2; 3; 1; 2; 3; 1; 2; 3;  
 3; 1; 2; 3; 1; 2; 3; 1; 2; 3; 1; 2; 3; 1; 2; 3; 1; 2; 3; 1;  
 1; 2; 3; 1; 2; 3; 1; 2; 3; 1; 2; 3; 1; 2; 3; 1; 2; 3; 1; 2;  
 2; 3; 1; 2; 3; 1; 2; 3; 1; 2; 3; 1; 2; 3; 1; 2; 3; 1; 2; 3;  
 3; 1; 2; 3; 1; 2; 3; 1; 2; 3; 1; 2; 3; 1; 2; 3; 1; 2; 3; 1;  
 1; 2; 3; 1; 2; 3; 1; 2; 3; 1; 2; 3; 1; 2; 3; 1; 2; 3; 1; 2;  
 2; 3; 1; 2; 3; 1; 2; 3; 1; 2; 3; 1; 2; 3; 1; 2; 3; 1; 2; 3;  
 3; 1; 2; 3; 1; 2; 3; 1; 2; 3; 1; 2; 3; 1; 2; 3; 1; 2; 3; 1;  
 1; 2; 3; 1; 2; 3; 1; 2; 3; 1; 2; 3; 1; 2; 3; 1; 2; 3; 1; 2;  
 2; 3; 1; 2; 3; 1; 2; 3; 1; 2; 3; 1; 2; 3; 1; 2; 3; 1; 2; 3;  
 3; 1; 2; 3; 1; 2; 3; 1; 2; 3; 1; 2; 3; 1; 2; 3; 1; 2; 3; 1;  
 1; 2; 3; 1; 2; 3; 1; 2; 3; 1; 2; 3; 1; 2; 3; 1; 2; 3; 1; 2;  
 2; 3; 1; 2; 3; 1; 2; 3; 1; 2; 3; 1; 2; 3; 1; 2; 3; 1; 2; 3;  
 3; 1; 2; 3; 1; 2; 3; 1; 2; 3; 1; 2; 3; 1; 2; 3; 1; 2; 3; 1;  
 ...]
```

This approach is quite limited, however. General-purpose cyclic data structures require mutation.[0 comments](#)

Modifying the List

Now, we'll start considering operations that mutate the list, starting with `insert_first`, which inserts an element at the front of the list:[0 comments](#)

```
let insert_first t value =
  let new_elt = { prev = None; next = !t; value } in
  begin match !t with
  | Some old_first -> old_first.prev <- Some new_elt
  | None -> ()
  end;
  t := Some new_elt;
  new_elt
```

OCaml * imperative-programming/dlist.ml , continued (part 3) * all code

`insert_first` first defines a new element `new_elt`, and then links it into the list, finally setting the list itself to point to `new_elt`. Note that the precedence of a `match` expression is very low, so to separate it from the following assignment (`t := Some new_elt`), we surround the match with `begin ... end`. We could have used parentheses for the same purpose. Without some kind of bracketing, the final assignment would incorrectly become part of the `None` case.[0 comments](#)

We can use `insert_after` to insert elements later in the list. `insert_after` takes as arguments both an element after which to insert the new node and a value to insert:[0 comments](#)

```
let insert_after elt value =
  let new_elt = { value; prev = Some elt; next = elt.next } in
  begin match elt.next with
  | Some old_next -> old_next.prev <- Some new_elt
  | None -> ()
  end;
  elt.next <- Some new_elt;
  new_elt
```

OCaml * imperative-programming/dlist.ml , continued (part 4) * all code

Finally, we need a `remove` function:[0 comments](#)

```
let remove t elt =
  let { prev; next; _ } = elt in
  begin match prev with
  | Some prev -> prev.next <- next
  | None -> t := next
  end;
  begin match next with
  | Some next -> next.prev <- prev;
  | None -> ()
  end;
  elt.prev <- None;
  elt.next <- None
```

OCaml * imperative-programming/dlist.ml , continued (part 5) * all code

Note that the preceding code is careful to change the `prev` pointer of the following element and the `next` pointer of the previous element, if they exist. If there's no previous element, then the list pointer itself is updated. In any case, the next and previous pointers of the element itself are set to `None`.[0 comments](#)

These functions are more fragile than they may seem. In particular, misuse of the interface may lead to corrupted data. For example, double-removing an element will cause the main list reference to be set to `None`, thus emptying the list. Similar problems arise from removing an element from a list it doesn't belong to.[0 comments](#)

This shouldn't be a big surprise. Complex imperative data structures can be quite tricky, considerably trickier than their pure equivalents. The issues described previously can be dealt with by more careful error detection, and such error correction is taken care of in modules like Core's `Doubly_linked`. You should use imperative data structures from a well-designed library when you can. And when you can't, you should make sure to put great care into your error handling.[0 comments](#)

Iteration Functions

When defining containers like lists, dictionaries, and trees, you'll typically want to define a set of iteration functions like `iter`, `map`, and `fold`, which let you concisely express common iteration patterns.[0 comments](#)

`Dlist` has two such iterators: `iter`, the goal of which is to call a unit-producing function on every element of the list, in order; and `find_el`, which runs a provided test function on each values stored in the list, returning the first element that passes the test. Both `iter` and `find_el` are implemented using simple recursive loops that use `next` to walk from element to element and `value` to extract the element from a given node:[0 comments](#)

```
let iter t ~f =
  let rec loop = function
    | None -> ()
    | Some el -> f (value el); loop (next el)
  in
  loop !t

let find_el t ~f =
  let rec loop = function
    | None -> None
    | Some elt ->
      if f (value elt) then Some elt
      else loop (next elt)
  in
  loop !t
```

OCaml * imperative-programming/dlist.ml , continued (part 6) * all code

This completes our implementation, but there's still considerably more work to be done to make a really usable doubly linked list. As mentioned earlier, you're probably better off using something like Core's `Doubly_linked` module that has a more complete interface and has more of the tricky corner cases worked out. Nonetheless, this example should serve to demonstrate some of the techniques you can use to build nontrivial imperative data structure in OCaml, as well as some of the pitfalls.[0 comments](#)

LAZINESS AND OTHER BENIGN EFFECTS

There are many instances where you basically want to program in a pure style, but you want to make limited use of side effects to improve the performance of your code. Such side effects are sometimes called *benign effects*, and they are a useful way of leveraging OCaml's imperative features while still maintaining most of the benefits of pure programming.[0 comments](#)

One of the simplest benign effects is *laziness*. A lazy value is one that is not computed until it is actually needed. In OCaml, lazy values are created using the `lazy` keyword, which can be used to convert any expression of type `s` into a lazy value of type `s Lazy.t`. The evaluation of that expression is delayed until forced with `Lazy.force`:[0 comments](#)

```
# let v = lazy (print_string "performing lazy computation\n"; sqrt 16.);;
val v : float lazy_t = <lazy>
# Lazy.force v;;

performing lazy computation
- : float = 4.
# Lazy.force v;;
- : float = 4.
```

OCaml Utop * imperative-programming/lazy.topscrip , continued (part 1) * all code

You can see from the `print` statement that the actual computation was performed only once, and only after `force` had been called.[0 comments](#)

To better understand how laziness works, let's walk through the implementation of our own lazy type. We'll start by declaring types to represent a lazy value:[0 comments](#)

```
# type 'a lazy_state =
  | Delayed of (unit -> 'a)
  | Value of 'a
  | Exn of exn
;;
type 'a lazy_state = Delayed of (unit -> 'a) | Value of 'a | Exn of exn
```

OCaml Utop * imperative-programming/lazy.topscrip , continued (part 2) * all code

A `lazy_state` represents the possible states of a lazy value. A lazy value is `Delayed` before it has been run, where `Delayed` holds a function for computing the value in question. A lazy value is in the `Value` state when it has been forced and the computation ended normally. The `Exn` case is for when the lazy value has been forced, but the computation ended with an exception. A lazy value is simply a `ref` containing a `lazy_state`, where the `ref` makes it possible to change from being in the `Delayed` state to being in the `Value` or `Exn` states.[0 comments](#)

We can create a lazy value from a thunk, i.e., a function that takes a unit argument. Wrapping an expression in a thunk is another way to suspend the computation of an expression:[0 comments](#)

```
# let create_lazy f = ref (Delayed f);;
val create_lazy : (unit -> 'a) -> 'a lazy_state ref = <fun>
# let v = create_lazy
  (fun () -> print_string "performing lazy computation\n"; sqrt 16.);;
val v : float lazy_state ref = {contents = Delayed <fun>}
```

OCaml Utop * imperative-programming/lazy.topscript , continued (part 3) * all code

Now we just need a way to force a lazy value. The following function does just that:[0 comments](#)

```
# let force v =
  match !v with
  | Value x -> x
  | Exn e -> raise e
  | Delayed f ->
    try
      let x = f () in
      v := Value x;
      x
    with exn ->
      v := Exn exn;
      raise exn
;;
val force : 'a lazy_state ref -> 'a = <fun>
```

OCaml Utop * imperative-programming/lazy.topscript , continued (part 4) * all code

Which we can use in the same way we used `Lazy.force`:[0 comments](#)

```
# force v;;

performing lazy computation
- : float = 4.
# force v;;
- : float = 4.
```

OCaml Utop * imperative-programming/lazy.topscript , continued (part 5) * all code

The main user-visible difference between our implementation of laziness and the built-in version is syntax. Rather than writing `create_lazy (fun () -> sqrt 16.)`, we can (with the built-in `lazy`) just write `lazy (sqrt 16.)`.[0 comments](#)

Memoization and Dynamic Programming

Another benign effect is *memoization*. A memoized function remembers the result of previous invocations of the function so that they can be returned without further computation when the same arguments are presented again.[0 comments](#)

Here's a function that takes as an argument an arbitrary single-argument function and returns a memoized version of that function. Here we'll use Core's `Hashtbl` module, rather than our toy `Dictionary`:[0 comments](#)

```
# let memoize f =
  let table = Hashtbl.Poly.create () in
  (fun x ->
    match Hashtbl.find table x with
    | Some y -> y
    | None ->
      let y = f x in
      Hashtbl.add_exn table ~key:x ~data:y;
      y
  );;
val memoize : ('a -> 'b) -> 'a -> 'b = <fun>
```

OCaml Utop * imperative-programming/memo.topscript , continued (part 1) * all code

The preceding code is a bit tricky. `memoize` takes as its argument a function `f` and then allocates a hash table (called `table`) and returns a new function as the memoized version of `f`. When called, this new function looks in `table` first, and if it fails to find a value, calls `f` and stashes the result in `table`. Note that `table` doesn't go out of scope as long as the function returned by `memoize` is in scope.[0 comments](#)

Memoization can be useful whenever you have a function that is expensive to recompute and you don't mind caching old values indefinitely. One important caution: a memoized function by its nature leaks memory. As long as you hold on to the memoized function, you're holding every result it has returned thus far.[0 comments](#)

Memoization is also useful for efficiently implementing some recursive algorithms. One good example is the algorithm for computing the *edit distance* (also called the Levenshtein distance) between two strings. The edit distance is the number of single-character changes (including letter switches, insertions, and deletions) required to convert one string to the other. This kind of distance metric can be useful for a variety of approximate string-matching problems, like spellcheckers.[0 comments](#)

Consider the following code for computing the edit distance. Understanding the algorithm isn't important here, but you should pay attention to the structure of the recursive calls:[0 comments](#)

```
# let rec edit_distance s t =
  match String.length s, String.length t with
  | (0,x) | (x,0) -> x
  | (len_s,len_t) ->
    let s' = String.drop_suffix s 1 in
    let t' = String.drop_suffix t 1 in
    let cost_to_drop_both =
      if s.[len_s - 1] = t.[len_t - 1] then 0 else 1
    in
    List.reduce_exn ~f:~f.Int.min
      [ edit_distance s' t + 1
        ; edit_distance s t' + 1
        ; edit_distance s t' + cost_to_drop_both
      ]
;;
val edit_distance : string -> string -> int = <fun>
# edit_distance "OCaml" "ocaml";;
- : int = 2
```

OCaml Utop * imperative-programming/memo.topscript , continued (part 2) * all code

The thing to note is that if you call `edit_distance "OCaml" "ocaml"`, then that will in turn dispatch the following calls:[0 comments](#)

```
edit_distance "OCam" "ocaml"
edit_distance "OCaml" "ocam"
edit_distance "OCam" "ocam"
```

Diagram * imperative-programming/edit_distance.ascii * all code

And these calls will in turn dispatch other calls:[0 comments](#)

```
edit_distance "OCam" "ocaml"
  edit_distance "OCa" "ocaml"
    edit_distance "OCam" "ocam"
      edit_distance "OCa" "ocam"
        edit_distance "OCaml" "ocam"
          edit_distance "OCam" "oca"
            edit_distance "OCaml" "oca"
              edit_distance "OCam" "oca"
                edit_distance "OCam" "oca"
                  edit_distance "OCa" "oca"
                    edit_distance "OCa" "oca"
```

Diagram * imperative-programming/edit_distance2.ascii * all code

As you can see, some of these calls are repeats. For example, there are two different calls to `edit_distance "OCam" "oca"`. The number of redundant calls grows exponentially with the size of the strings, meaning that our implementation of `edit_distance` is brutally slow for large strings. We can see this by writing a small timing function:[0 comments](#)

```
# let time f =
  let start = Time.now () in
```

```

let x = f () in
let stop = Time.now () in
printf "Time: %s\n" (Time.Span.to_string (Time.diff stop start));
x ;;
val time : (unit -> 'a) -> 'a = <fun>

```

OCaml Utop * imperative-programming/memo.topscript , continued (part 3) * all code

And now we can use this to try out some examples:[0 comments](#)

```

# time (fun () -> edit_distance "OCaml" "ocaml");;

Time: 0.698805ms
- : int = 2
# time (fun () -> edit_distance "OCaml 4.01" "ocaml 4.01");;

Time: 1.78467s
- : int = 2

```

OCaml Utop * imperative-programming/memo.topscript , continued (part 4) * all code

Just those few extra characters made it thousands of times slower![0 comments](#)

Memoization would be a huge help here, but to fix the problem, we need to memoize the calls that `edit_distance` makes to itself. This technique is sometimes referred to as *dynamic programming*. To see how to do this, let's step away from `edit_distance` and instead consider a much simpler example: computing the n th element of the Fibonacci sequence. The Fibonacci sequence by definition starts out with two 1s, with every subsequent element being the sum of the previous two. The classic recursive definition of Fibonacci is as follows:[0 comments](#)

```

# let rec fib i =
  if i <= 1 then 1 else fib (i - 1) + fib (i - 2);;
val fib : int -> int = <fun>

```

OCaml Utop * imperative-programming/fib.topscript , continued (part 1) * all code

This is, however, exponentially slow, for the same reason that `edit_distance` was slow: we end up making many redundant calls to `fib`. It shows up quite dramatically in the performance:[0 comments](#)

```

# time (fun () -> fib 20);;

Time: 0.379086ms
- : int = 10946
# time (fun () -> fib 40);;

Time: 4.61983s
- : int = 165580141

```

OCaml Utop * imperative-programming/fib.topscript , continued (part 2) * all code

As you can see, `fib 40` takes thousands of times longer to compute than `fib 20`.[0 comments](#)

So, how can we use memoization to make this faster? The tricky bit is that we need to insert the memoization before the recursive calls within `fib`. We can't just define `fib` in the ordinary way and memoize it after the fact and expect the first call to `fib` to be improved:[0 comments](#)

```

# let fib = memoize fib;;
val fib : int -> int = <fun>
# time (fun () -> fib 40);;

Time: 4.90749s
- : int = 165580141
# time (fun () -> fib 40);;

Time: 0.00286102ms
- : int = 165580141

```

OCaml Utop * imperative-programming/fib.topscript , continued (part 3) * all code

In order to make `fib` fast, our first step will be to rewrite `fib` in a way that unwinds the

recursion. The following version expects as its first argument a function (called `fib`) that will be called in lieu of the usual recursive call:[0 comments](#)

```
# let fib_norec fib i =
  if i <= 1 then i
  else fib (i - 1) + fib (i - 2) ;;
val fib_norec : (int -> int) -> int -> int = <fun>
```

OCaml Utop * imperative-programming/fib.topscript , continued (part 4) * all code

We can now turn this back into an ordinary Fibonacci function by tying the recursive knot:[0 comments](#)

```
# let rec fib i = fib_norec fib i;;
val fib : int -> int = <fun>
# fib 20;;
- : int = 6765
```

OCaml Utop * imperative-programming/fib.topscript , continued (part 5) * all code

We can even write a polymorphic function that we'll call `make_rec` that can tie the recursive knot for any function of this form:[0 comments](#)

```
# let make_rec f_norec =
  let rec f x = f_norec f x in
  f
;;
val make_rec : (('a -> 'b) -> 'a -> 'b) -> 'a -> 'b = <fun>
# let fib = make_rec fib_norec;;
val fib : int -> int = <fun>
# fib 20;;
- : int = 6765
```

OCaml Utop * imperative-programming/fib.topscript , continued (part 6) * all code

This is a pretty strange piece of code, and it may take a few moments of thought to figure out what's going on. Like `fib_norec`, the function `f_norec` passed into `make_rec` is a function that isn't recursive but takes as an argument a function that it will call. What `make_rec` does is to essentially feed `f_norec` to itself, thus making it a true recursive function.[0 comments](#)

This is clever enough, but all we've really done is find a new way to implement the same old slow Fibonacci function. To make it faster, we need a variant of `make_rec` that inserts memoization when it ties the recursive knot. We'll call that function `memo_rec`:[0 comments](#)

```
# let memo_rec f_norec x =
  let fref = ref (fun _ -> assert false) in
  let f = memoize (fun x -> f_norec !fref x) in
  fref := f;
  f x
;;
val memo_rec : (('a -> 'b) -> 'a -> 'b) -> 'a -> 'b = <fun>
```

OCaml Utop * imperative-programming/fib.topscript , continued (part 7) * all code

Note that `memo_rec` has the same signature as `make_rec`.[0 comments](#)

We're using the reference here as a way of tying the recursive knot without using a `let rec`, which for reasons we'll describe later wouldn't work here.[0 comments](#)

Using `memo_rec`, we can now build an efficient version of `fib`:[0 comments](#)

```
# let fib = memo_rec fib_norec;;
val fib : int -> int = <fun>
# time (fun () -> fib 40);;

Time: 0.0371933ms
- : int = 102334155
```

OCaml Utop * imperative-programming/fib.topscript , continued (part 8) * all code

And as you can see, the exponential time complexity is now gone.[0 comments](#)

The memory behavior here is important. If you look back at the definition of `memo_rec`, you'll see that the call `memo_rec fib_norec` does not trigger a call to `memoize`. Only when `fib` is called and thereby the final argument to `memo_rec` is presented does `memoize` get called. The result of that

We can use `memo_rec` as part of a single declaration that makes this look like it's little more than a special form of `let`: `let rec:0 comments`

OCaml Utop * imperative-programming/fib.topscript , continued (part 9) * all code

But memoization is a good approach for optimizing `edit_distance`, and we can apply the same approach we used on `fib` here. We will need to change `edit_distance` to take a pair of strings as a single argument, since `memo_rec` only works on single-argument functions. (We can always recover the original interface with a wrapper function.) With just that change and the addition of the `memo_rec` call, we can get a memoized version of `edit_distance`: [0 comments](#)

OCaml Utop * imperative-programming/memo.topscript , continued (part 6) * all code

OCaml Utop * imperative-programming/memo.topscript , continued (part 7) * all code

OCaml * imperative-programming/let_rec.ml * all code

Note that x is an ordinary value, not a function. As such, it's not clear how this

definition should be handled by the compiler. You could imagine it compiling down to an infinite loop, but `x` is of type `int`, and there's no `int` that corresponds to an infinite loop. As such, this construct is effectively impossible to compile.[0 comments](#)

To avoid such impossible cases, the compiler only allows three possible constructs to show up on the righthand side of a `let rec`: a function definition, a constructor, or the `lazy` keyword. This excludes some reasonable things, like our definition of `memo_rec`, but it also blocks things that don't make sense, like our definition of `x`.[0 comments](#)

It's worth noting that these restrictions don't show up in a lazy language like Haskell. Indeed, we can make something like our definition of `x` work if we use OCaml's laziness:[0 comments](#)

```
# let rec x = lazy (Lazy.force x + 1);;
val x : int lazy_t = <lazy>
```

OCaml Utop * imperative-programming/letrec.topscrip , continued (part 2) * all code

Of course, actually trying to compute this will fail. OCaml's `lazy` throws an exception when a lazy value tries to force itself as part of its own evaluation.[0 comments](#)

```
# Lazy.force x;;
Exception: Lazy.Undefined.
```

OCaml Utop * imperative-programming/letrec.topscrip , continued (part 3) * all code

But we can also create useful recursive definitions with `lazy`. In particular, we can use laziness to make our definition of `memo_rec` work without explicit mutation.[0 comments](#)

```
# let lazy_memo_rec f_norec x =
  let rec f = lazy (memoize (fun x -> f_norec (Lazy.force f) x)) in
  (Lazy.force f) x
;;
val lazy_memo_rec : (('a -> 'b) -> 'a -> 'b) -> 'a -> 'b = <fun>
# time (fun () -> lazy_memo_rec fib_norec 40);;
```

```
Time: 0.041008ms
- : int = 102334155
```

OCaml Utop * imperative-programming/letrec.topscrip , continued (part 5) * all code

Laziness is more constrained than explicit mutation, and so in some cases can lead to code whose behavior is easier to think about.[0 comments](#)

INPUT AND OUTPUT

Imperative programming is about more than modifying in-memory data structures. Any function that doesn't boil down to a deterministic transformation from its arguments to its return value is imperative in nature. That includes not only things that mutate your program's data, but also operations that interact with the world outside of your program. An important example of this kind of interaction is I/O, i.e., operations for reading or writing data to things like files, terminal input and output, and network sockets.[0 comments](#)

There are multiple I/O libraries in OCaml. In this section we'll discuss OCaml's buffered I/O library that can be used through the `In_channel` and `Out_channel` modules in `Core`. Other I/O primitives are also available through the `Unix` module in `Core` as well as `Async`, the asynchronous I/O library that is covered in [Chapter 18, Concurrent Programming with Async](#). Most of the functionality in `Core`'s `In_channel` and `Out_channel` (and in `Core`'s `Unix` module) derives from the standard library, but we'll use `Core`'s interfaces here.[0 comments](#)

Terminal I/O

OCaml's buffered I/O library is organized around two types: `in_channel`, for channels you read from, and `out_channel`, for channels you write to. The `In_channel` and `Out_channel` modules only have direct support for channels corresponding to files and terminals; other kinds of channels can be created through the `Unix` module.[0 comments](#)

We'll start our discussion of I/O by focusing on the terminal. Following the UNIX model, communication with the terminal is organized around three channels, which correspond to the three standard file descriptors in Unix:[0 comments](#)

`In_channel.stdin`

The "standard input" channel. By default, input comes from the terminal, which handles keyboard input.[0 comments](#)

`Out_channel.stdout`

The "standard output" channel. By default, output written to `stdout` appears on the user terminal.[0 comments](#)

`Out_channel.stderr`

The "standard error" channel. This is similar to `stdout` but is intended for error messages.[0 comments](#)

The values `stdin`, `stdout`, and `stderr` are useful enough that they are also available in the global namespace directly, without having to go through the `In_channel` and `Out_channel` modules.[0 comments](#)

Let's see this in action in a simple interactive application. The following program, `time_converter`, prompts the user for a time zone, and then prints out the current time in that time zone. Here, we use Core's `Zone` module for looking up a time zone, and the `Time` module for computing the current time and printing it out in the time zone in question:[0 comments](#)

```
open Core.Std

let () =
  Out_channel.output_string stdout "Pick a timezone: ";
  Out_channel.flush stdout;
  match In_channel.input_line stdin with
  | None -> failwith "No timezone provided"
  | Some zone_string ->
    let zone = Zone.find_exn zone_string in
    let time_string = Time.to_string_abs (Time.now ()) ~zone in
    Out_channel.output_string stdout
      (String.concat
        ["The time in "; Zone.to_string zone; " is "; time_string; ".\n"]);
    Out_channel.flush stdout
```

OCaml * imperative-programming/time_converter.ml * all code

We can build this program using **corebuild** and run it. You'll see that it prompts you for input, as follows:[0 comments](#)

```
$ corebuild time_converter.byte
$ ./time_converter.byte
Pick a timezone:
```

Terminal * imperative-programming/time_converter.out * all code

You can then type in the name of a time zone and hit Return, and it will print out the current time in the time zone in question:[0 comments](#)

```
Pick a timezone: Europe/London
The time in Europe/London is 2013-08-15 00:03:10.666220+01:00.
```

Terminal * imperative-programming/time_converter2.out * all code

We called `Out_channel.flush` on `stdout` because `out_channels` are buffered, which is to say that OCaml doesn't immediately do a write every time you call `output_string`. Instead, writes are buffered until either enough has been written to trigger the flushing of the buffers, or until a flush is explicitly requested. This greatly increases the efficiency of the writing process by reducing the number of system calls.[0 comments](#)

Note that `In_channel.input_line` returns a string option, with `None` indicating that the input stream has ended (i.e., an end-of-file condition). `Out_channel.output_string` is used to print the final output, and `Out_channel.flush` is called to flush that output to the screen. The final flush is not technically required, since the program ends after that instruction, at which point all remaining output will be flushed anyway, but the explicit flush is nonetheless good practice.[0 comments](#)

Formatted Output with `printf`

Generating output with functions like `Out_channel.output_string` is simple and easy to

understand, but can be a bit verbose. OCaml also supports formatted output using the `printf` function, which is modeled after `printf` in the C standard library. `printf` takes a *format string* that describes what to print and how to format it, as well as arguments to be printed, as determined by the formatting directives embedded in the format string. So, for example, we can write:

```
# printf "%i is an integer, %F is a float, \"%s\" is a string\n"
  3 4.5 "five";;
```

```
3 is an integer, 4.5 is a float, "five" is a string
- : unit = ()
```

OCaml Utop * imperative-programming/printf.topscript , continued (part 1) * all code

Unlike C's `printf`, the `printf` in OCaml is type-safe. In particular, if we provide an argument whose type doesn't match what's presented in the format string, we'll get a type error:

```
# printf "An integer: %i\n" 4.5;;
```

Characters 26-29:

```
Error: This expression has type float but an expression was expected of type
      int
```

OCaml Utop * imperative-programming/printf.topscript , continued (part 2) * all code

Understanding Format Strings

The format strings used by `printf` turn out to be quite different from ordinary strings. This difference ties to the fact that OCaml format strings, unlike their equivalent in C, are type-safe. In particular, the compiler checks that the types referred to by the format string match the types of the rest of the arguments passed to `printf`.

To check this, OCaml needs to analyze the contents of the format string at compile time, which means the format string needs to be available as a string literal at compile time. Indeed, if you try to pass an ordinary string to `printf`, the compiler will complain:

```
# let fmt = "%i is an integer, %F is a float, \"%s\" is a string\n";;
val fmt : string = "%i is an integer, %F is a float, \"%s\" is a string\n"
# printf fmt 3 4.5 "five";;
Characters 9-12:
Error: This expression has type string but an expression was expected of type
      ('a -> 'b -> 'c -> 'd, out_channel, unit) format =
        ('a -> 'b -> 'c -> 'd, out_channel, unit, unit, unit, unit)
        format6
```

OCaml Utop * imperative-programming/printf.topscript , continued (part 3) * all code

If OCaml infers that a given string literal is a format string, then it parses it at compile time as such, choosing its type in accordance with the formatting directives it finds. Thus, if we add a type annotation indicating that the string we're defining is actually a format string, it will be interpreted as such:

```
# let fmt : ('a, 'b, 'c) format =
  "%i is an integer, %F is a float, \"%s\" is a string\n";;
val fmt : (int -> float -> string -> 'c, 'b, 'c) format = <abstr>
```

OCaml Utop * imperative-programming/printf.topscript , continued (part 4) * all code

And accordingly, we can pass it to `printf`:

```
# printf fmt 3 4.5 "five";;
```

```
3 is an integer, 4.5 is a float, "five" is a string
- : unit = ()
```

OCaml Utop * imperative-programming/printf.topscript , continued (part 5) * all code

If this looks different from everything else you've seen so far, that's because it is. This is really a special case in the type system. Most of the time, you don't need to worry about this special handling of format strings—you can just use `printf` and not worry about the details. But it's useful to keep the broad outlines of the story in the back of your head.

Now let's see how we can rewrite our time conversion program to be a little more concise using `printf`:

```

open Core.Std

let () =
  printf "Pick a timezone: %!";
  match In_channel.input_line stdin with
  | None -> failwith "No timezone provided"
  | Some zone_string ->
    let zone = Time.Zone.find_exn zone_string in
    let time_string = Time.to_string_abs (Time.now ()) ~zone in
    printf "The time in %s is %s.\n%" (Time.Zone.to_string zone) time_string

```

OCaml * imperative-programming/time_converter2.ml * all code

In the preceding example, we've used only two formatting directives: %s, for including a string, and %! which causes printf to flush the channel.[0 comments](#)

printf's formatting directives offer a significant amount of control, allowing you to specify things like:[0 comments](#)

- Alignment and padding[0 comments](#)
- Escaping rules for strings[0 comments](#)
- Whether numbers should be formatted in decimal, hex, or binary[0 comments](#)
- Precision of float conversions[0 comments](#)

There are also printf-style functions that target outputs other than stdout, including:[0 comments](#)

- eprintf, which prints to stderr[0 comments](#)
- fprintf, which prints to an arbitrary out_channel[0 comments](#)
- sprintf, which returns a formatted string[0 comments](#)

All of this, and a good deal more, is described in the API documentation for the Printf module in the OCaml Manual.[0 comments](#)

File I/O

Another common use of in_channels and out_channels is for working with files. Here are a couple of functions—one that creates a file full of numbers, and the other that reads in such a file and returns the sum of those numbers:[0 comments](#)

```

# let create_number_file filename numbers =
  let outc = Out_channel.create filename in
  List.iter numbers ~f:(fun x -> fprintf outc "%d\n" x);
  Out_channel.close outc
;;
val create_number_file : string -> int list -> unit = <fun>
# let sum_file filename =
  let file = In_channel.create filename in
  let numbers = List.map ~f:Int.of_string (In_channel.input_lines file) in
  let sum = List.fold ~init:0 ~f:(+) numbers in
  In_channel.close file;
  sum
;;
val sum_file : string -> int = <fun>
# create_number_file "numbers.txt" [1;2;3;4;5];;
- : unit = ()
# sum_file "numbers.txt";;
- : int = 15

```

OCaml Utop * imperative-programming/file.topscript , continued (part 1) * all code

For both of these functions, we followed the same basic sequence: we first create the channel, then use the channel, and finally close the channel. The closing of the channel is important, since without it, we won't release resources associated with the file back to the operating system.[0 comments](#)

One problem with the preceding code is that if it throws an exception in the middle of its work, it won't actually close the file. If we try to read a file that doesn't actually contain numbers, we'll see such an error:[0 comments](#)

```

# sum_file "/etc/hosts";;
Exception: (Failure "Int.of_string: \"##\"").

```

And if we do this over and over in a loop, we'll eventually run out of file descriptors:[0 comments](#)

```
# for i = 1 to 10000 do try ignore (sum_file "/etc/hosts") with _ -> () done;;
- : unit = ()
# sum_file "numbers.txt";;
Exception: (Sys_error "numbers.txt: Too many open files").
```

OCaml Utop * imperative-programming/file.topscript , continued (part 3) * all code

And now, you'll need to restart your toplevel if you want to open any more files![0 comments](#)

To avoid this, we need to make sure that our code cleans up after itself. We can do this using the `protect` function described in [Chapter 7, Error Handling](#), as follows:[0 comments](#)

```
# let sum_file filename =
  let file = In_channel.create filename in
  protect ~f:(fun () ->
    let numbers = List.map ~f:Int.of_string (In_channel.input_lines file) in
    List.fold ~init:0 ~f:(+) numbers)
    ~finally:(fun () -> In_channel.close file)
  ;;
val sum_file : string -> int = <fun>
```

OCaml Utop * imperative-programming/file2.topscript , continued (part 1) * all code

And now, the file descriptor leak is gone:[0 comments](#)

```
# for i = 1 to 10000 do try ignore (sum_file "/etc/hosts") with _ -> () done;;
- : unit = ()
# sum_file "numbers.txt";;
- : int = 15
```

OCaml Utop * imperative-programming/file2.topscript , continued (part 2) * all code

This is really an example of a more general issue with imperative programming. When programming imperatively, you need to be quite careful to make sure that exceptions don't leave you in an awkward state.[0 comments](#)

`In_channel` has functions that automate the handling of some of these details. For example, `In_channel.with_file` takes a filename and a function for processing data from an `in_channel` and takes care of the bookkeeping associated with opening and closing the file. We can rewrite `sum_file` using this function, as shown here:[0 comments](#)

```
# let sum_file filename =
  In_channel.with_file filename ~f:(fun file ->
    let numbers = List.map ~f:Int.of_string (In_channel.input_lines file) in
    List.fold ~init:0 ~f:(+) numbers)
  ;;
val sum_file : string -> int = <fun>
```

OCaml Utop * imperative-programming/file2.topscript , continued (part 3) * all code

Another misfeature of our implementation of `sum_file` is that we read the entire file into memory before processing it. For a large file, it's more efficient to process a line at a time. You can use the `In_channel.fold_lines` function to do just that:[0 comments](#)

```
# let sum_file filename =
  In_channel.with_file filename ~f:(fun file ->
    In_channel.fold_lines file ~init:0 ~f:(fun sum line ->
      sum + Int.of_string line))
  ;;
val sum_file : string -> int = <fun>
```

OCaml Utop * imperative-programming/file2.topscript , continued (part 4) * all code

This is just a taste of the functionality of `In_channel` and `Out_channel`. To get a fuller understanding, you should review the API documentation for those modules.[0 comments](#)

ORDER OF EVALUATION

The order in which expressions are evaluated is an important part of the definition of a programming language, and it is particularly important when programming imperatively. Most programming languages you're likely to have encountered are *strict*, and OCaml is, too. In a strict

language, when you bind an identifier to the result of some expression, the expression is evaluated before the variable is bound. Similarly, if you call a function on a set of arguments, those arguments are evaluated before they are passed to the function.[0 comments](#)

Consider the following simple example. Here, we have a collection of angles, and we want to determine if any of them have a negative `sin`. The following snippet of code would answer that question:[0 comments](#)

```
# let x = sin 120. in
  let y = sin 75. in
  let z = sin 128. in
  List.exists ~f:(fun x -> x < 0.) [x;y;z]
;;
- : bool = true
```

OCaml Utop * imperative-programming/order.topscript , continued (part 1) * all code

In some sense, we don't really need to compute the `sin 128.` because `sin 75.` is negative, so we could know the answer before even computing `sin 128.`[0 comments](#)

It doesn't have to be this way. Using the `lazy` keyword, we can write the original computation so that `sin 128.` won't ever be computed:[0 comments](#)

```
# let x = lazy (sin 120.) in
  let y = lazy (sin 75.) in
  let z = lazy (sin 128.) in
  List.exists ~f:(fun x -> Lazy.force x < 0.) [x;y;z]
;;
- : bool = true
```

OCaml Utop * imperative-programming/order.topscript , continued (part 2) * all code

We can confirm that fact by a few well-placed `printf`s:[0 comments](#)

```
# let x = lazy (printf "1\n"; sin 120.) in
  let y = lazy (printf "2\n"; sin 75.) in
  let z = lazy (printf "3\n"; sin 128.) in
  List.exists ~f:(fun x -> Lazy.force x < 0.) [x;y;z]
;;

1
2
- : bool = true
```

OCaml Utop * imperative-programming/order.topscript , continued (part 3) * all code

OCaml is strict by default for a good reason: lazy evaluation and imperative programming generally don't mix well because laziness makes it harder to reason about when a given side effect is going to occur. Understanding the order of side effects is essential to reasoning about the behavior of an imperative program.[0 comments](#)

In a strict language, we know that expressions that are bound by a sequence of `let` bindings will be evaluated in the order that they're defined. But what about the evaluation order within a single expression? Officially, the answer is that evaluation order within an expression is undefined. In practice, OCaml has only one compiler, and that behavior is a kind of de facto standard. Unfortunately, the evaluation order in this case is often the opposite of what one might expect.[0 comments](#)

Consider the following example:[0 comments](#)

```
# List.exists ~f:(fun x -> x < 0.)
[ (printf "1\n"; sin 120.);
  (printf "2\n"; sin 75.);
  (printf "3\n"; sin 128.); ]
;;

3
2
1
- : bool = true
```

OCaml Utop * imperative-programming/order.topscript , continued (part 4) * all code

Here, you can see that the subexpression that came last was actually evaluated first! This is generally the case for many different kinds of expressions. If you want to make sure of the

evaluation order of different subexpressions, you should express them as a series of `let` bindings.[0 comments](#)

SIDE EFFECTS AND WEAK POLYMORPHISM

Consider the following simple, imperative function:[0 comments](#)

```
# let remember =
  let cache = ref None in
  (fun x ->
    match !cache with
    | Some y -> y
    | None -> cache := Some x; x)
;;
val remember : '_a -> '_a = <fun>
```

OCaml Utop * imperative-programming/weak.topscript , continued (part 1) * all code

`remember` simply caches the first value that's passed to it, returning that value on every call. That's because `cache` is created and initialized once and is shared across invocations of `remember`.[0 comments](#)

`remember` is not a terribly useful function, but it raises an interesting question: what is its type?[0 comments](#)

On its first call, `remember` returns the same value it's passed, which means its input type and return type should match. Accordingly, `remember` should have type `t -> t` for some type `t`. There's nothing about `remember` that ties the choice of `t` to any particular type, so you might expect OCaml to generalize, replacing `t` with a polymorphic type variable. It's this kind of generalization that gives us polymorphic types in the first place. The identity function, as an example, gets a polymorphic type in this way:[0 comments](#)

```
# let identity x = x;;
val identity : 'a -> 'a = <fun>
# identity 3;;
- : int = 3
# identity "five";;
- : string = "five"
```

OCaml Utop * imperative-programming/weak.topscript , continued (part 2) * all code

As you can see, the polymorphic type of `identity` lets it operate on values with different types.[0 comments](#)

This is not what happens with `remember`, though. As you can see from the above examples, the type that OCaml infers for `remember` looks almost, but not quite, like the type of the identity function. Here it is again:[0 comments](#)

```
val remember : '_a -> '_a = <fun>
```

OCaml * imperative-programming/remember_type.ml * all code

The underscore in the type variable `'_a` tells us that the variable is only *weakly polymorphic*, which is to say that it can be used with any *single* type. That makes sense because, unlike `identity`, `remember` always returns the value it was passed on its first invocation, which means its return value must always have the same type.[0 comments](#)

OCaml will convert a weakly polymorphic variable to a concrete type as soon as it gets a clue as to what concrete type it is to be used as:[0 comments](#)

```
# let remember_three () = remember 3;;
val remember_three : unit -> int = <fun>
# remember;;
- : int -> int = <fun>
# remember "avocado";;
Characters 9-18:
Error: This expression has type string but an expression was expected of type
      int
```

OCaml Utop * imperative-programming/weak.topscript , continued (part 3) * all code

Note that the type of `remember` was settled by the definition of `remember_three`, even though `remember_three` was never called![0 comments](#)

The Value Restriction

So, when does the compiler infer weakly polymorphic types? As we've seen, we need weakly polymorphic types when a value of unknown type is stored in a persistent mutable cell. Because the type system isn't precise enough to determine all cases where this might happen, OCaml uses a rough rule to flag cases that don't introduce any persistent mutable cells, and to only infer polymorphic types in those cases. This rule is called *the value restriction*.[0 comments](#)

The core of the value restriction is the observation that some kinds of expressions, which we'll refer to as *simple values*, by their nature can't introduce persistent mutable cells, including:[0 comments](#)

- Constants (i.e., things like integer and floating-point literals)[0 comments](#)
- Constructors that only contain other simple values[0 comments](#)
- Function declarations, i.e., expressions that begin with `fun` or `function`, or the equivalent `let` binding, `let f x = ...`[0 comments](#)
- `let` bindings of the form `let var = expr1 in expr2`, where both `expr1` and `expr2` are simple values[0 comments](#)

Thus, the following expression is a simple value, and as a result, the types of values contained within it are allowed to be polymorphic:[0 comments](#)

```
# (fun x -> [x;x]);;  
- : 'a -> 'a list = <fun>
```

OCaml Utop * imperative-programming/value_restriction.topscript , continued (part 1) * all code

But, if we write down an expression that isn't a simple value by the preceding definition, we'll get different results. For example, consider what happens if we try to memoize the function defined previously.[0 comments](#)

```
# memoize (fun x -> [x;x]);;  
- : '_a -> '_a list = <fun>
```

OCaml Utop * imperative-programming/value_restriction.topscript , continued (part 2) * all code

The memoized version of the function does in fact need to be restricted to a single type because it uses mutable state behind the scenes to cache values returned by previous invocations of the function. But OCaml would make the same determination even if the function in question did no such thing. Consider this example:[0 comments](#)

```
# identity (fun x -> [x;x]);;  
- : '_a -> '_a list = <fun>
```

OCaml Utop * imperative-programming/value_restriction.topscript , continued (part 3) * all code

It would be safe to infer a fully polymorphic variable here, but because OCaml's type system doesn't distinguish between pure and impure functions, it can't separate those two cases.[0 comments](#)

The value restriction doesn't require that there is no mutable state, only that there is no *persistent* mutable state that could share values between uses of the same function. Thus, a function that produces a fresh reference every time it's called can have a fully polymorphic type:[0 comments](#)

```
# let f () = ref None;;  
val f : unit -> 'a option ref = <fun>
```

OCaml Utop * imperative-programming/value_restriction.topscript , continued (part 4) * all code

But a function that has a mutable cache that persists across calls, like `memoize`, can only be weakly polymorphic.[0 comments](#)

Partial Application and the Value Restriction

Most of the time, when the value restriction kicks in, it's for a good reason, i.e., it's because the value in question can actually only safely be used with a single type. But sometimes, the value restriction kicks in when you don't want it. The most common such case is partially applied functions. A partially applied function, like any function application, is not a simple value, and as such, functions created by partial application are sometimes less general than you might

[expect.0 comments](#)

Consider the `List.init` function, which is used for creating lists where each element is created by calling a function on the index of that element:[0 comments](#)

```
# List.init;;
- : int -> f:(int -> 'a) -> 'a list = <fun>
# List.init 10 ~f:Int.to_string;;
- : string list = ["0"; "1"; "2"; "3"; "4"; "5"; "6"; "7"; "8"; "9"]
```

OCaml Utop * imperative-programming/value_restriction.topscript , continued (part 5) * all code

Imagine we wanted to create a specialized version of `List.init` that always created lists of length 10. We could do that using partial application, as follows:[0 comments](#)

```
# let list_init_10 = List.init 10;;
val list_init_10 : f:(int -> 'a) -> 'a list = <fun>
```

OCaml Utop * imperative-programming/value_restriction.topscript , continued (part 6) * all code

As you can see, we now infer a weakly polymorphic type for the resulting function. That's because there's nothing that guarantees that `List.init` isn't creating a persistent ref somewhere inside of it that would be shared across multiple calls to `list_init_10`. We can eliminate this possibility, and at the same time get the compiler to infer a polymorphic type, by avoiding partial application:[0 comments](#)

```
# let list_init_10 ~f = List.init 10 ~f;;
val list_init_10 : f:(int -> 'a) -> 'a list = <fun>
```

OCaml Utop * imperative-programming/value_restriction.topscript , continued (part 7) * all code

This transformation is referred to as *eta expansion* and is often useful to resolve problems that arise from the value restriction.[0 comments](#)

Relaxing the Value Restriction

OCaml is actually a little better at inferring polymorphic types than was suggested previously. The value restriction as we described it is basically a syntactic check: you can do a few operations that count as simple values, and anything that's a simple value can be generalized.[0 comments](#)

But OCaml actually has a relaxed version of the value restriction that can make use of type information to allow polymorphic types for things that are not simple values.[0 comments](#)

For example, we saw that a function application, even a simple application of the identity function, is not a simple value and thus can turn a polymorphic value into a weakly polymorphic one:[0 comments](#)

```
# identity (fun x -> [x;x]);;
- : '_a -> '_a list = <fun>
```

OCaml Utop * imperative-programming/value_restriction.topscript , continued (part 8) * all code

But that's not always the case. When the type of the returned value is immutable, then OCaml can typically infer a fully polymorphic type:[0 comments](#)

```
# identity [];;
- : 'a list = []
```

OCaml Utop * imperative-programming/value_restriction.topscript , continued (part 9) * all code

On the other hand, if the returned type is potentially mutable, then the result will be weakly polymorphic:[0 comments](#)

```
# [[]];;
- : 'a array = [[]]
# identity [[]];;
- : '_a array = [[]]
```

OCaml Utop * imperative-programming/value_restriction.topscript , continued (part 10) * all code

A more important example of this comes up when defining abstract data types. Consider the following simple data structure for an immutable list type that supports constant-time concatenation:[0 comments](#)

```

# module Concat_list : sig
  type 'a t
  val empty : 'a t
  val singleton : 'a -> 'a t
  val concat : 'a t -> 'a t -> 'a t (* constant time *)
  val to_list : 'a t -> 'a list (* linear time *)
end = struct

  type 'a t = Empty | Singleton of 'a | Concat of 'a t * 'a t

  let empty = Empty
  let singleton x = Singleton x
  let concat x y = Concat (x,y)

  let rec to_list_with_tail t tail =
    match t with
    | Empty -> tail
    | Singleton x -> x :: tail
    | Concat (x,y) -> to_list_with_tail x (to_list_with_tail y tail)

  let to_list t =
    to_list_with_tail t []

end;;
module Concat_List :
  sig
    type 'a t
    val empty : 'a t
    val singleton : 'a -> 'a t
    val concat : 'a t -> 'a t -> 'a t
    val to_list : 'a t -> 'a list
  end

```

OCaml Utop * imperative-programming/value_restriction.topscript , continued (part 11) * all code

The details of the implementation don't matter so much, but it's important to note that a `Concat_list.t` is unquestionably an immutable value. However, when it comes to the value restriction, OCaml treats it as if it were mutable:[0 comments](#)

```

# Concat_list.empty;;
- : 'a Concat_list.t = <abstr>
# identity Concat_list.empty;;
- : '_a Concat_list.t = <abstr>

```

OCaml Utop * imperative-programming/value_restriction.topscript , continued (part 12) * all code

The issue here is that the signature, by virtue of being abstract, has obscured the fact that `Concat_list.t` is in fact an immutable data type. We can resolve this in one of two ways: either by making the type concrete (i.e., exposing the implementation in the `mli`), which is often not desirable; or by marking the type variable in question as *covariant*. We'll learn more about covariance and contravariance in [Chapter 11, Objects](#), but for now, you can think of it as an annotation that can be put in the interface of a pure data structure.[0 comments](#)

In particular, if we replace `type 'a t` in the interface with `type +'a t`, that will make it explicit in the interface that the data structure doesn't contain any persistent references to values of type `'a`, at which point, OCaml can infer polymorphic types for expressions of this type that are not simple values:[0 comments](#)

```

# module Concat_list : sig
  type +'a t
  val empty : 'a t
  val singleton : 'a -> 'a t
  val concat : 'a t -> 'a t -> 'a t (* constant time *)
  val to_list : 'a t -> 'a list (* linear time *)
end = struct

  type 'a t = Empty | Singleton of 'a | Concat of 'a t * 'a t

  ...

end;;
module Concat_List :
  sig
    type +'a t
    val empty : 'a t
    val singleton : 'a -> 'a t
    val concat : 'a t -> 'a t -> 'a t
    val to_list : 'a t -> 'a list
  end

```


Now, we can apply the identity function to `Concat_list.empty` without without losing any polymorphism:[0 comments](#)

```
# identity Concat_list.empty;;  
- : 'a Concat_list.t = <abstr>
```

OCaml Utop * imperative-programming/value_restriction.topscript , continued (part 14) * all code

SUMMARY

This chapter has covered quite a lot of ground, including:[0 comments](#)

- Discussing the building blocks of mutable data structures as well as the basic imperative constructs like `for` loops, `while` loops, and the sequencing operator `;` [0 comments](#)
- Walking through the implementation of a couple of classic imperative data structures [0 comments](#)
- Discussing so-called benign effects like memoization and laziness [0 comments](#)
- Covering OCaml's API for blocking I/O [0 comments](#)
- Discussing how language-level issues like order of evaluation and weak polymorphism interact with OCaml's imperative features [0 comments](#)

The scope and sophistication of the material here is an indication of the importance of OCaml's imperative features. The fact that OCaml defaults to immutability shouldn't obscure the fact that imperative programming is a fundamental part of building any serious application, and that if you want to be an effective OCaml programmer, you need to understand OCaml's approach to imperative programming. [0 comments](#)

[< Previous](#)[Next >](#)